

LLOV: A Fast Static Data-Race Checker for OpenMP Programs



Utpal Bora* Santanu Das Pankaj Kukreja Saurabh Joshi
Ramakrishna Upadrasta Sanjay Rajopadhye

*PhD Student
Computer Science and Engineering
IIT Hyderabad, India

January 18, 2021

- Missing data sharing clauses

```
1 #pragma omp parallel for \  
2     private (temp,i,j)  
3     for (i = 0; i < len; i++)  
4         for (j = 0; j < len; j++){  
5             temp = u[i][j];  
6             sum = sum + temp * temp;  
7         }
```

Listing 1: DRB021: OpenMP Worksharing construct with data race

- Missing data sharing clauses
- Loop carried dependences

```
1  for (i=0;i<n;i++) {  
2  #pragma omp parallel for  
3      for (j=1;j<m;j++) {  
4      b[i][j] = b[i][j-1] ;  
5      }  
6  }
```

Listing 2: DRB038: Example with Loop Carried Dependence

- Missing data sharing clauses
- Loop carried dependences
- SIMD races

```
1 #pragma omp simd
2 for (int i=0; i<len-1; i++){
3     a[i+1] = a[i] + b[i];
4 }
```

Listing 3: DRB024: Example with SIMD data race

- Missing data sharing clauses
- Loop carried dependences
- SIMD races
- Synchronization issues

```
1 #pragma omp parallel \  
2     shared(b, error)  
3     {  
4 #pragma omp for nowait  
5     for(i = 0; i < len; i++)  
6         a[i] = b + a[i]*5;  
7 #pragma omp single  
8     error = a[9] + 1;  
9 }
```

Listing 4: DRB013: Example with data race due to improper synchronization

- Missing data sharing clauses
- Loop carried dependences
- SIMD races
- Synchronization issues
- Control flow dependent on number of threads

```
1 #pragma omp parallel
2   if (omp_get_thread_num() % 2
3       == 0) {
4       Flag = true;
5   }
```

Listing 5: Control flow dependent on number of threads

LLOV is a **language agnostic**, **static** OpenMP data race checker in the LLVM compiler framework.

LLOV is a **language agnostic**, **static** OpenMP data race checker in the LLVM compiler framework. LLOV

- is based on intermediate representation of LLVM (**LLVM-IR**)

LLOV is a **language agnostic**, **static** OpenMP data race checker in the LLVM compiler framework. LLOV

- is based on intermediate representation of LLVM (**LLVM-IR**)
- can handle FORTRAN as well as C/C++

LLOV is a **language agnostic**, **static** OpenMP data race checker in the LLVM compiler framework. LLOV

- is based on intermediate representation of LLVM (**LLVM-IR**)
- can handle FORTRAN as well as C/C++
- uses Polyhedral framework, **Polly**, of LLVM

LLOV is a **language agnostic**, **static** OpenMP data race checker in the LLVM compiler framework. LLOV

- is based on intermediate representation of LLVM (**LLVM-IR**)
- can handle FORTRAN as well as C/C++
- uses Polyhedral framework, **Polly**, of LLVM
- can conservatively state when a program is **data race free**

LLOV is a **language agnostic**, **static** OpenMP data race checker in the LLVM compiler framework. LLOV

- is based on intermediate representation of LLVM (**LLVM-IR**)
- can handle FORTRAN as well as C/C++
- uses Polyhedral framework, **Polly**, of LLVM
- can conservatively state when a program is **data race free**
- is capable of generating task graphs of OpenMP constructs

LLOV: Flow Diagram

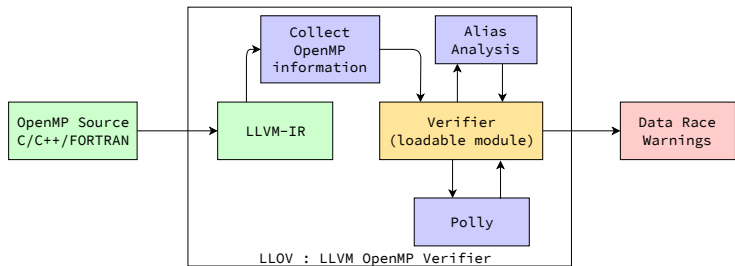


Figure 1: Flow Diagram of LLVM OpenMP Verifier (LLOV)

```
1 #pragma omp parallel shared(b, error)
2 {
3     #pragma omp for nowait
4     for(i = 0; i < len; i++)
5         a[i] = b + a[i]*5;
6     #pragma omp single
7         error = a[9] + 1;
8 }
```

```
1 #pragma omp parallel shared(b, error)
2 {
3     #pragma omp for nowait
4         for(i = 0; i < len; i++)
5             a[i] = b + a[i]*5;
6     #pragma omp single
7         error = a[9] + 1;
8 }
```

Directive: OMP_Parallel

Variables:

Private: %.omp.ub = alloca i32, align 4

Private: %.omp.lb = alloca i32, align 4

Shared: i32* %i

Shared: i32* %len

Firstprivate: i64 %vla

Shared: i32* %a

Shared: i32* %b

Shared: i32* %error

Private: %.omp.stride = alloca i32, align 4

Private: %.omp.is_last = alloca i32, align 4

Child Directives:

1: Directive: OMP_Workshare_Loop

Schedule type : Static Schedule (auto-chunked)

2: Directive: OMP_Workshare_single

3: Directive: OMP_Barrier

Listing 7: In-memory representation of a directive

```

1 #pragma omp parallel shared(b, error)
2 {
3 #pragma omp for nowait
4     for(i = 0; i < len; i++)
5         a[i] = b + a[i]*5;
6 #pragma omp single
7     error = a[9] + 1;
8 }

```

Directive: OMP_Parallel

Variables:

```

Private:  %.omp.ub = <Directive> ::= <Dtype> [ Sched ] { <Var> } { <Directive> }
Private:  %.omp.lb = <Dtype> ::= parallel | for | simd
Shared:  i32* %i      | workshare | single
Shared:  i32* %len    | master | critical
Firstprivate: i64 %via <Var> ::= <Vtype> val
Shared:  i32* %a      <Vtype> ::= private | firstprivate
Shared:  i32* %b      | shared | lastprivate
Shared:  i32* %error
Private:  %.omp.stride = alloca i32, allg reduction | threadprivate
Private:  %.omp.is_lia<Sched>oca i32::=a[i<modifier> ] [ ordered ] <Stype> <chunk>

```

Child Directives:

```

1: Directive:  OMP_For<Stype>loop ::= static | dynamic | guided | auto | runtime
   Schedule type : Static<chunk>e (auto::=positive-int-const
2: Directive:  OMP_Workshare single
3: Directive:  OMP_Barrier

```

Listing 7: In-memory representation of a directive


```
1 for (i=0; i<m ; i++) {  
2 #pragma omp parallel for  
3     for (j=1; j<n ;j++) {  
4 SO:    b[i][j] = b[i][j-1];  
5     }  
6 }
```

Iteration Domain : $\mathbf{I} = \{\text{SO}(i, j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $\mathbf{S} = \{\text{SO}(i, j) \rightarrow (i, j)\} \cap_{dom} \mathbf{I}$

Access Map : $\mathbf{A} = \{\text{SO}(i, j) \rightarrow \mathbf{M}(i, j); \text{SO}(i, j) \rightarrow \mathbf{M}(i, j-1)\}$

Dependences : $\mathbf{D} = \{\text{SO}(i, j) \rightarrow (i, j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i=0; i<m ; i++) {  
2 #pragma omp parallel for  
3     for (j=1; j<n ;j++) {  
4 SO:    b[i][j] = b[i][j-1];  
5     }  
6 }
```

Iteration Domain : $\mathbf{I} = \{\text{SO}(i,j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $\mathbf{S} = \{\text{SO}(i,j) \rightarrow (i,j)\} \cap_{dom} \mathbf{I}$

Access Map : $\mathbf{A} = \{\text{SO}(i,j) \rightarrow \mathbf{M}(i,j); \text{SO}(i,j) \rightarrow \mathbf{M}(i,j-1)\}$

Dependences : $\mathbf{D} = \{\text{SO}(i,j) \rightarrow (i,j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i=0; i<m ; i++) {  
2 #pragma omp parallel for  
3     for (j=1; j<n ;j++) {  
4 S0:   b[i][j] = b[i][j-1];  
5     }  
6 }
```

Iteration Domain : $\mathbf{I} = \{S0(i,j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $\mathbf{S} = \{S0(i,j) \rightarrow (i,j)\} \cap_{dom} \mathbf{I}$

Access Map : $\mathbf{A} = \{S0(i,j) \rightarrow M(i,j); S0(i,j) \rightarrow M(i,j-1)\}$

Dependences : $\mathbf{D} = \{S0(i,j) \rightarrow (i,j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i=0; i<m ; i++) {  
2 #pragma omp parallel for  
3   for (j=1; j<n ;j++) {  
4 SO:   b[i][j] = b[i][j-1];  
5   }  
6 }
```

Iteration Domain : $\mathbf{I} = \{\text{SO}(i, j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $\mathbf{S} = \{\text{SO}(i, j) \rightarrow (i, j)\} \cap_{\text{dom}} \mathbf{I}$

Access Map : $\mathbf{A} = \{\text{SO}(i, j) \rightarrow \mathbf{M}(i, j); \text{SO}(i, j) \rightarrow \mathbf{M}(i, j-1)\}$

Dependences : $\mathbf{D} = \{\text{SO}(i, j) \rightarrow (i, j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i=0; i<m ; i++) {  
2 #pragma omp parallel for  
3   for (j=1; j<n ; j++) {  
4 S0:   b[i][j] = b[i][j-1];  
5   }  
6 }
```

Iteration Domain : $\mathbf{I} = \{S0(i, j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $\mathbf{S} = \{S0(i, j) \rightarrow (i, j)\} \cap_{dom} \mathbf{I}$

Access Map : $\mathbf{A} = \{S0(i, j) \rightarrow M(i, j); S0(i, j) \rightarrow M(i, j-1)\}$

Dependences : $\mathbf{D} = \{S0(i, j) \rightarrow (i, j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i=0; i<m ; i++) {  
2 #pragma omp parallel for  
3   for (j=1; j<n ;j++) {  
4 S0:  b[i][j] = b[i][j-1];  
5   }  
6 }
```

Iteration Domain : $\mathbf{I} = \{S0(i,j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $\mathbf{S} = \{S0(i,j) \rightarrow (i,j)\} \cap_{dom} \mathbf{I}$

Access Map : $\mathbf{A} = \{S0(i,j) \rightarrow M(i,j); S0(i,j) \rightarrow M(i,j-1)\}$

Dependences : $\mathbf{D} = \{S0(i,j) \rightarrow (i,j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i=0; i<m ; i++) {  
2 #pragma omp parallel for  
3     for (j=1; j<n ;j++) {  
4 S0:   b[i][j] = b[i][j-1];  
5     }  
6 }
```

Iteration Domain : $\mathbf{I} = \{S0(i,j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $\mathbf{S} = \{S0(i,j) \rightarrow (i,j)\} \cap_{dom} \mathbf{I}$

Access Map : $\mathbf{A} = \{S0(i,j) \rightarrow M(i,j); S0(i,j) \rightarrow M(i,j-1)\}$

Dependences : $\mathbf{D} = \{S0(i,j) \rightarrow (i,j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i=0; i<m ; i++) {  
2 #pragma omp parallel for  
3   for (j=1; j<n ;j++) {  
4 SO:  b[i][j] = b[i][j-1];  
5   }  
6 }
```

Iteration Domain : $\mathbf{I} = \{\text{SO}(i,j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $\mathbf{S} = \{\text{SO}(i,j) \rightarrow (i,j)\} \cap_{dom} \mathbf{I}$

Access Map : $\mathbf{A} = \{\text{SO}(i,j) \rightarrow \mathbf{M}(i,j); \text{SO}(i,j) \rightarrow \mathbf{M}(i,j-1)\}$

Dependences : $\mathbf{D} = \{\text{SO}(i,j) \rightarrow (i,j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$


```
1 for (i=0; i<m ; i++) {  
2 #pragma omp parallel for  
3     for (j=1; j<n ;j++) {  
4 SO:   b[i][j] = b[i][j-1];  
5     }  
6 }
```

Iteration Domain : $\mathbf{I} = \{\text{SO}(i,j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $\mathbf{S} = \{\text{SO}(i,j) \rightarrow (i,j)\} \cap_{dom} \mathbf{I}$

Access Map : $\mathbf{A} = \{\text{SO}(i,j) \rightarrow \mathbf{M}(i,j); \text{SO}(i,j) \rightarrow \mathbf{M}(i,j-1)\}$

Dependences : $\mathbf{D} = \{\text{SO}(i,j) \rightarrow (i,j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1  for (i=0;i<10;i++) {  
2  #pragma omp parallel for  
3      for (j=1;j<10;j++) {  
4          b[i][j]=b[i][j-1];  
5      }  
6  }
```

Listing 10: Example with Loop Carried Dependence

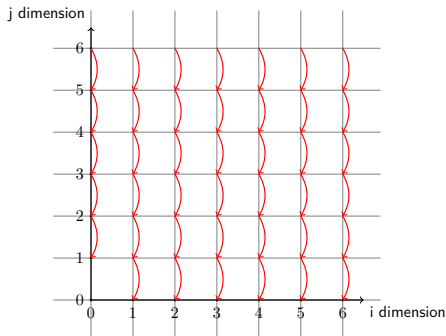


Figure 2: Dependence Polyhedra

```
1  for (i=0;i<10;i++) {  
2  #pragma omp parallel for  
3      for (j=1;j<10;j++) {  
4          b[i][j]=b[i][j-1];  
5      }  
6  }
```

Listing 11: Example with Loop Carried Dependence

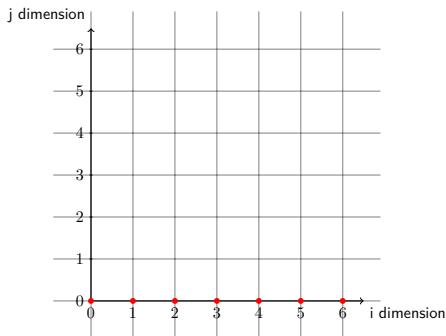


Figure 3: Projection of the Dependence Polyhedra on i-dimension

Zero magnitude of the projections on a dimension signifies that the dimension is **parallel**.

```
1  for (i=0;i<10;i++) {  
2  #pragma omp parallel for  
3      for (j=1;j<10;j++) {  
4          b[i][j]=b[i][j-1];  
5      }  
6  }
```

Listing 12: Example with Loop Carried Dependence

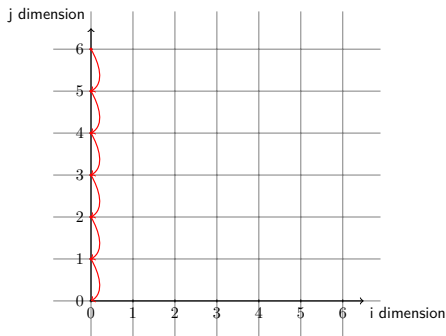


Figure 4: Projection of the Dependence Polyhedra on j-dimension

Non-zero magnitude of the projections on a dimension signifies that the dimension is **not parallel**.

Benchmarks:

- DataRaceBench C/C++ v1.2 [Liao et al., 2018a, Liao et al., 2018b]
- OmpSCR v2.0 [Dorta et al., 2004, Dorta et al., 2005]
- DataRaceBench FORTRAN [Kukreja et al., 2019]

Benchmarks:

- DataRaceBench C/C++ v1.2 [Liao et al., 2018a, Liao et al., 2018b]
- OmpSCR v2.0 [Dorta et al., 2004, Dorta et al., 2005]
- DataRaceBench FORTRAN [Kukreja et al., 2019]

System Specifications:

System: Two Intel Xeon E5-2697 v4 @ 2.30GHz processors

OS: 64 bit Ubuntu 18.04.2 LTS server

Kernel: Linux kernel version 4.15.0-48-generic

Threads: 72 (2 x 36) hardware threads

Memory: 128GB

OpenMP library: LLVM OpenMP runtime v5.0.1 (libomp5)

Table 1: Race detection tools with the version numbers used for comparison

Tools	Source	Version / Commit
HELGRIND [Valgrind-project, 2007b]	Valgrind	3.13.0
VALGRIND DRD [Valgrind-project, 2007a]	Valgrind	3.13.0
TSAN-LLVM [Serebryany and Iskhodzhanov, 2009]	LLVM	6.0.1
ARCHER [Atzeni et al., 2016]	git master branch	fc17353
SWORD [Atzeni et al., 2018]	git master branch	7a08f3c

Results: DataRaceBench v1.2 comparison

Table 2: Maximum number of Races reported by different tools in DataRaceBench 1.2

Tools	Race: Yes		Race: No		Coverage/116
	TP	FN	TN	FP	
HELGRIND	56	3	2	55	116
VALGRIND DRD	56	3	26	31	116
TSAN-LLVM	57	2	2	55	116
ARCHER	56	3	2	55	116
SWORD	47	4	24	4	79
LLOV	48	2	36	5	91

Results: DataRaceBench v1.2 comparison

Table 2: Maximum number of Races reported by different tools in DataRaceBench 1.2

Tools	Race: Yes		Race: No		Coverage/116
	TP	FN	TN	FP	
HELGRIND	56	3	2	55	116
VALGRIND DRD	56	3	26	31	116
TSAN-LLVM	57	2	2	55	116
ARCHER	56	3	2	55	116
SWORD	47	4	24	4	79
LLOV	48	2	36	5	91

Table 3: Maximum number of Races reported by different tools in common 61 kernels of DataRaceBench 1.2

Tools	Race: Yes		Race: No		Coverage/61
	TP	FN	TN	FP	
HELGRIND	42	1	2	16	61
VALGRIND DRD	42	1	12	6	61
TSAN-LLVM	42	1	2	16	61
ARCHER	42	1	2	16	61
SWORD	42	1	17	1	61
LLOV	42	1	16	2	61

Table 4: Performance of the tools on DataRaceBench 1.2

Tools	Precision	Recall	Accuracy	F1 Score	Diagnostic odds ratio
HELGRIND	0.50	0.95	0.50	0.66	0.68
VALGRIND DRD	0.64	0.95	0.71	0.77	15.66
TSAN-LLVM	0.51	0.97	0.51	0.67	1.04
ARCHER	0.50	0.95	0.50	0.66	0.68
SWORD	0.92	0.92	0.90	0.92	70.50
LLOV	0.91	0.96	0.92	0.93	172.80

Table 4: Performance of the tools on DataRaceBench 1.2

Tools	Precision	Recall	Accuracy	F1 Score	Diagnostic odds ratio
HELGRIND	0.50	0.95	0.50	0.66	0.68
VALGRIND DRD	0.64	0.95	0.71	0.77	15.66
TSAN-LLVM	0.51	0.97	0.51	0.67	1.04
ARCHER	0.50	0.95	0.50	0.66	0.68
SWORD	0.92	0.92	0.90	0.92	70.50
LLOV	0.91	0.96	0.92	0.93	172.80

Table 5: Performance of the tools on **common 61** kernels of DataRaceBench 1.2

Tools	Precision	Recall	Accuracy	F1 Score	Diagnostic odds ratio
HELGRIND	0.72	0.98	0.72	0.83	5.25
VALGRIND DRD	0.88	0.98	0.89	0.92	84.00
TSAN-LLVM	0.72	0.98	0.72	0.83	5.25
ARCHER	0.72	0.98	0.72	0.83	5.25
SWORD	0.98	0.98	0.97	0.98	714.00
LLOV	0.95	0.98	0.95	0.97	336.00

Results: DataRaceBench v1.2 runtime

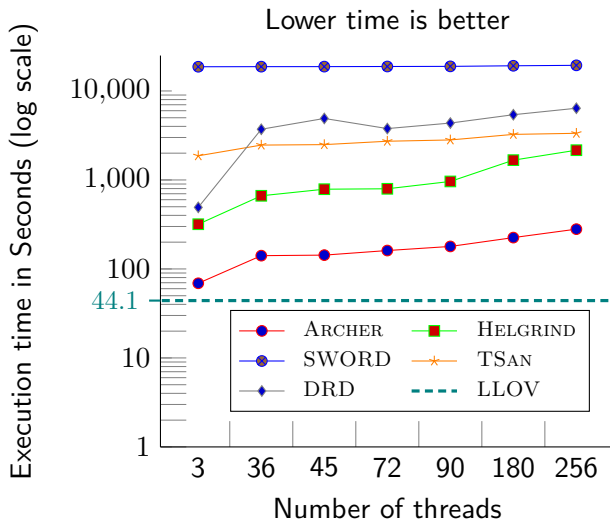


Figure 5: DataRaceBench v1.2 total execution time by different tools on logarithmic scale

Results: DataRaceBench v1.2 runtime

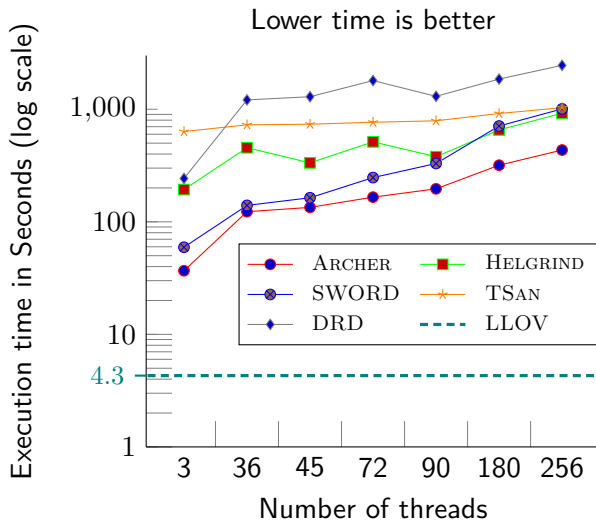


Figure 6: DataRaceBench v1.2 total time taken by different tools for common 61 kernels on logarithmic scale

Table 6: Comparison of different tools on OmpSCR v2.0

Tools	Race: Yes		Race: No		Coverage/14
	TP	FN	TN	FP	
HELGRIND	8	0	0	9	14
VALGRIND DRD	8	0	2	5	14
TSAN-LLVM	7	1	2	6	14
ARCHER	7	1	2	4	14
SWORD	3	4	3	0	10
LLOV	4	1	2	5	10

Results: OmpSCR v2.0 runtime

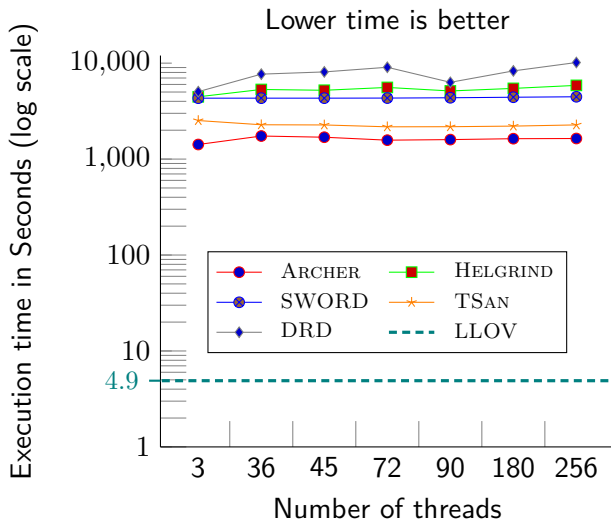


Figure 7: OmpSCR v2.0 total execution time by different tools on logarithmic scale

An implementation of DataRaceBench C/C++ v1.2 [Liao et al., 2018b] in FORTRAN 95.

- Converted 92 (out of 116) C/C++ kernels to FORTRAN
- Demonstrate that LLOV is language agnostic
- Already open-sourced this benchmark [Kukreja et al., 2019]

Table 7: Maximum number of Races reported by different tools in DataRaceBench FORTRAN

Tools	Race: Yes		Race: No		Coverage/92
	TP	FN	TN	FP	
HELGRIND	46	6	4	36	92
VALGRIND DRD	45	7	21	19	92
LLOV	36	7	19	5	67

LLOV: A Fast Static Data-Race Checker for OpenMP Programs



LLOV is freely available for download.

Link: <https://github.com/utpalbora/llov>

Blog: <https://compilers.cse.iith.ac.in/projects/llov/>

Open source links:

- DataRaceBench FORTRAN:
https://github.com/IITH-Compilers/dr_b_fortran
- LLOV source: Please drop me an email at
cs14mtech11017@iith.ac.in

Contributions Welcome!!

We welcome your contributions in any form. **Thank You!**



Atzeni, S., Gopalakrishnan, G., Rakamaric, Z., Ahn, D. H., Laguna, I., Schulz, M., Lee, G. L., Protze, J., and Müller, M. S. (2016).

Archer: effectively spotting data races in large openmp applications.

In *Parallel and Distributed Processing Symposium, 2016 IEEE International*, pages 53–62, Chicago, IL, USA. IEEE, IEEE.



Atzeni, S., Gopalakrishnan, G., Rakamaric, Z., Laguna, I., Lee, G. L., and Ahn, D. H. (2018).

Sword: A bounded memory-overhead detector of openmp data races in production runs.

In *2018 IEEE International Parallel and Distributed Processing Symposium (IPDPS)*, pages 845–854, Vancouver, BC, Canada. IEEE, IEEE.



Dorta, A. J., Rodriguez, C., and de Sande, F. (2004).

OpenMP Source Code Repository.

<https://sourceforge.net/projects/ompscr/files/OmpSCR/>.

[Online; accessed 19-May-2019].



Dorta, A. J., Rodriguez, C., and de Sande, F. (2005).

The openmp source code repository.

In *13th Euromicro Conference on Parallel, Distributed and Network-Based Processing*, pages 244–250, Washington, DC, USA. IEEE Computer Society.



Kukreja, P., Shukla, H., and Bora, U. (2019).
DataRaceBench FORTRAN.
https://github.com/IITH-Compilers/drb_fortran.
[Online; accessed 19-October-2019].



Liao, C., Lin, P.-H., Asplund, J., Schordan, M., and Karlin, I. (2018a).
DataRaceBench v1.2.0.
<https://github.com/LLNL/dataracebench>.
[Online; accessed 19-May-2019].



Liao, C., Lin, P.-H., Schordan, M., and Karlin, I. (2018b).
A semantics-driven approach to improving dataracebench's openmp standard coverage.
In *Evolving OpenMP for Evolving Architectures*, pages 189–202, Cham. Springer
International Publishing.



Serebryany, K. and Iskhodzhanov, T. (2009).
Threadsanitizer: Data race detection in practice.
In *Proceedings of the Workshop on Binary Instrumentation and Applications*, WBIA '09,
pages 62–71, New York, NY, USA. ACM.



Valgrind-project (2007a).
DRD: a thread error detector.
<http://valgrind.org/docs/manual/drd-manual.html>.
[Online; accessed 08-May-2019].



Valgrind-project (2007b).

Helgrind: a thread error detector.

<http://valgrind.org/docs/manual/hg-manual.html>.

[Online; accessed 08-May-2019].