

Data Race detection in OpenMP programs using the LLVM Framework



Utpal Bora

PhD Student
Computer Science and Engineering
IIT Hyderabad, India

April 3rd, 2021

- 1 Introduction to OpenMP
- 2 What are Data Races?
- 3 Static data race detection in LLVM
- 4 Results
- 5 Hands-on with LLOV

- 1 Introduction to OpenMP
- 2 What are Data Races?
- 3 Static data race detection in LLVM
- 4 Results
- 5 Hands-on with LLOV

Open Multi-Processing (OpenMP) is a framework for C, C++ and FORTRAN programs to enable **shared memory** parallel programming on a multi-threaded system.

Open Multi-Processing (OpenMP) is a framework for C, C++ and FORTRAN programs to enable **shared memory** parallel programming on a multi-threaded system.

OpenMP **Supports:**

- *Single Program Multiple Data* (SMPD) programming model

Open Multi-Processing (OpenMP) is a framework for C, C++ and FORTRAN programs to enable **shared memory** parallel programming on a multi-threaded system.

OpenMP **Supports:**

- *Single Program Multiple Data* (SMPD) programming model
- *Single Instruction Multiple Data* (SIMD) programming model

Open Multi-Processing (OpenMP) is a framework for C, C++ and FORTRAN programs to enable **shared memory** parallel programming on a multi-threaded system.

OpenMP **Supports:**

- *Single Program Multiple Data* (SMPD) programming model
- *Single Instruction Multiple Data* (SIMD) programming model
- Divide-and-conquer paradigm with tasks

Open Multi-Processing (OpenMP) is a framework for C, C++ and FORTRAN programs to enable **shared memory** parallel programming on a multi-threaded system.

OpenMP **Supports:**

- *Single Program Multiple Data* (SMPD) programming model
- *Single Instruction Multiple Data* (SIMD) programming model
- Divide-and-conquer paradigm with tasks
- Data environment for shared memory consistency

Open Multi-Processing (OpenMP) is a framework for C, C++ and FORTRAN programs to enable **shared memory** parallel programming on a multi-threaded system.

OpenMP **Supports:**

- *Single Program Multiple Data* (SMPD) programming model
- *Single Instruction Multiple Data* (SIMD) programming model
- Divide-and-conquer paradigm with tasks
- Data environment for shared memory consistency
- Mutual exclusion and atomicity

Open Multi-Processing (OpenMP) is a framework for C, C++ and FORTRAN programs to enable **shared memory** parallel programming on a multi-threaded system.

OpenMP **Supports:**

- *Single Program Multiple Data* (SMPD) programming model
- *Single Instruction Multiple Data* (SIMD) programming model
- Divide-and-conquer paradigm with tasks
- Data environment for shared memory consistency
- Mutual exclusion and atomicity
- Implicit and explicit synchronization

Open Multi-Processing (OpenMP) is a framework for C, C++ and FORTRAN programs to enable **shared memory** parallel programming on a multi-threaded system.

OpenMP **Supports:**

- *Single Program Multiple Data* (SMPD) programming model
- *Single Instruction Multiple Data* (SIMD) programming model
- Divide-and-conquer paradigm with tasks
- Data environment for shared memory consistency
- Mutual exclusion and atomicity
- Implicit and explicit synchronization

OpenMP is **portable** across architectures from different vendors. (e.g. CPU, GPU, and FPGA [Mayer et al., 2019])

OpenMP Example: `#pragma omp parallel`

```
1 #pragma omp parallel \  
2   num_threads(3)  
3 {  
4   for (int i = 0; i < 15; i++) {  
5     printf("Iteration Id %d by \  
6       Thread Id: %d\n", i,  
7       omp_get_thread_num());  
8   }  
9 }
```

Listing 1: OpenMP `parallel` construct

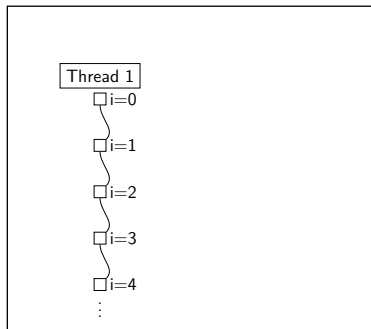


Figure 1: OpenMP thread execution model for `parallel` construct

OpenMP Example: `#pragma omp parallel`

```
1 #pragma omp parallel \  
2   num_threads(3)  
3 {  
4   for (int i = 0; i < 15; i++) {  
5     printf("Iteration Id %d by \  
6       Thread Id: %d\n", i,  
7       omp_get_thread_num());  
8   }  
9 }
```

Listing 2: OpenMP `parallel` construct

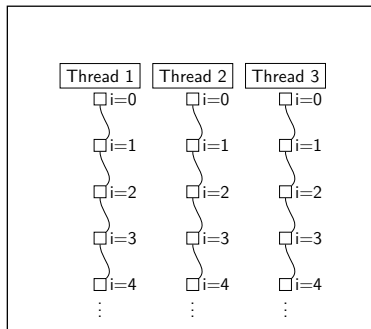


Figure 2: OpenMP thread execution model for `parallel` construct

OpenMP Example: `#pragma omp parallel for`



```
1 #pragma omp parallel for \  
2     schedule(static) \  
3     num_threads(3)  
4 for (int i = 0; i < 15; i++) {  
5     printf("Iteration Id %d by \  
6         Thread Id: %d\n", i,  
7         omp_get_thread_num());  
8 }
```

Listing 3: OpenMP worksharing `for` construct

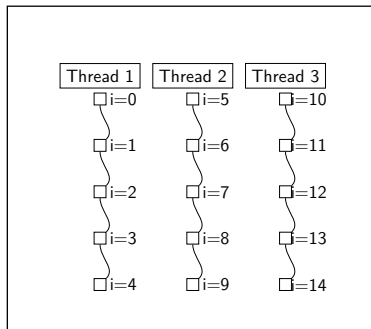


Figure 3: OpenMP thread execution model for `parallel for` construct

OpenMP Example: `#pragma omp parallel for` for



```
1 #pragma omp parallel for \  
2     schedule(static) \  
3     num_threads(3)  
4 for (int i = 0; i < 11; i++) {  
5     printf("Iteration Id %d by \  
6         Thread Id: %d\n", i,  
7         omp_get_thread_num());  
8 }
```

Listing 4: OpenMP worksharing `for` construct

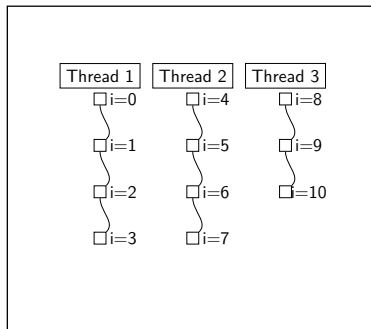


Figure 4: OpenMP thread execution model for `parallel for` construct

OpenMP Example: `#pragma omp parallel for` for



```
1 #pragma omp parallel for \  
2     schedule(static,1) \  
3     num_threads(3)  
4 for (int i = 0; i < 15; i++) {  
5     printf("Iteration Id %d by \  
6         Thread Id: %d\n", i,  
7         omp_get_thread_num());  
8 }
```

Listing 5: OpenMP worksharing `for` construct

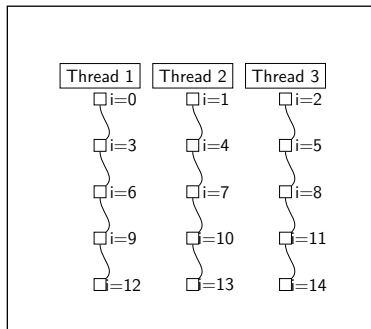


Figure 5: OpenMP thread execution model for `parallel for` construct

OpenMP Example: `#pragma omp parallel for` for



```
1 #pragma omp parallel for \  
2     schedule(static,4) \  
3     num_threads(3) \  
4 for (int i = 0; i < 15; i++) { \  
5     printf("Iteration Id %d by \  
6         Thread Id: %d\n", i, \  
7         omp_get_thread_num()); \  
8 }
```

Listing 6: OpenMP worksharing `for` construct

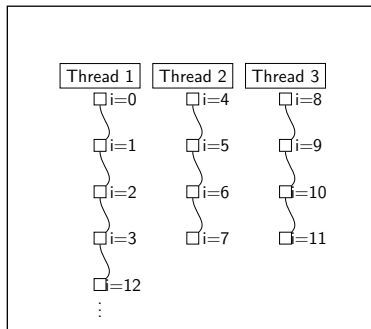


Figure 6: OpenMP thread execution model for `parallel for` construct with `chunking`

OpenMP Example: `#pragma omp parallel for` for



```
1 #pragma omp parallel for \  
2     schedule(static,4) \  
3     num_threads(3)  
4 for (int i = 0; i < 15; i++) {  
5     for (int j = 0; j < 15; j++) {  
6         printf("Iteration Id %d %d by \  
7             Thread Id: %d\n", i, j,  
8             omp_get_thread_num());  
9     }  
10 }
```

Listing 7: OpenMP worksharing `for` construct

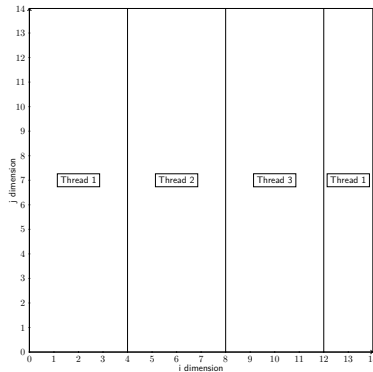


Figure 7: OpenMP thread execution model for `parallel for` construct with `chunking`

OpenMP Example: `#pragma omp parallel for` for



```
1 for (int i = 0; i < 15; i++) {  
2 #pragma omp parallel for \  
3     schedule(static,4) \  
4     num_threads(3)  
5     for (int j = 0; j < 15; j++) {  
6         printf("Iteration Id %d %d by \  
7             Thread Id: %d\n", i, j,  
8             omp_get_thread_num());  
9     }  
10 }
```

Listing 8: OpenMP worksharing `for` construct

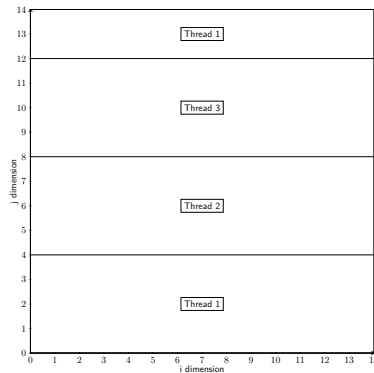


Figure 8: OpenMP thread execution model for `parallel for` construct with `chunking`

Table 1: Dependence Types

Dependence Type	RAW	WAR
Laxically Backward	<pre>for(int i=0;i<N;i++){ ... = A[i]; A[i+4] = ...; }</pre>	<pre>for(int i=0;i<N;i++){ A[i] = ...; ... = A[i+4]; }</pre>
Laxically Forward	<pre>for(int i=0;i<N;i++){ A[i+4] = ...; ... = A[i]; }</pre>	<pre>for(int i=0;i<N;i++){ ... = A[i]; A[i-4] = ...; }</pre>

Table 2: Vectorization and Parallelization legality

Dependence Type	Vectorization Legality	Parallelization Legality
Laxically Backward	Legal for $VF < \text{Iteration Distance}$	Restrictive schedule with $PF < \text{Iteration Distance}$
Laxically Forward	Legal with $VF = \infty$	Restrictive schedule with $PF < \text{Iteration Distance}$

Vectorization Example



```
1 for (int i = 0; i < 20; i++) {  
2     for (int j = 4; j < 20; j++) {  
3         b[i][j]=b[i][j-4] + x;  
4     }  
5 }
```

Listing 9: Example for Vectorization

Vectorization Example

```
1 for (int i = 0; i < 20; i++) {  
2     for (int j = 4; j < 20; j++) {  
3         b[i][j]=b[i][j-4] + x;  
4     }  
5 }
```

Listing 9: Example for Vectorization

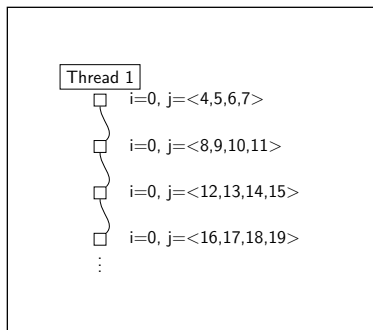


Figure 9: SIMD execution model within CPU

Parallelization Example:

```
1 for (int i = 0; i < 20; i++) {  
2 #pragma omp parallel for \  
3     schedule(static,4) \  
4     num_threads(3)  
5     for (int j = 4; j < 20; j++) {  
6         b[i][j] = b[i][j-4] + x;  
7     }  
8 }
```

Listing 10: OpenMP worksharing **for** construct
(incorrect parallelization)

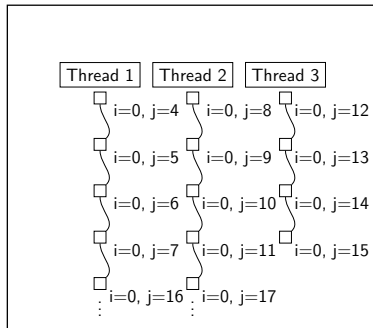


Figure 10: OpenMP thread execution schedule for
parallel for construct with data races

Parallelization Example:

```
1 for (int i = 0; i < 20; i++) {  
2 #pragma omp parallel for \  
3   schedule(static,4) \  
4   num_threads(3)  
5   for (int j = 4; j < 20; j++) {  
6     b[i][j] = b[i][j-4] + x;  
7 #pragma omp barrier  
8   }  
9 }
```

Listing 11: OpenMP worksharing **for** construct with barrier (incorrect parallelization)

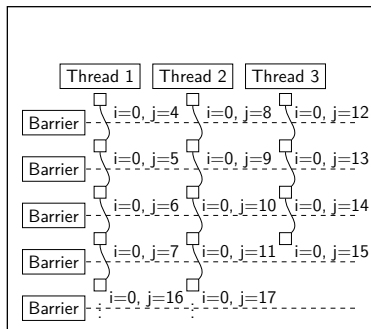


Figure 11: Possible OpenMP thread execution schedule for **parallel for** construct with **chunking**

Parallelization Example:

```
1 for (int i = 0; i < 20; i++) {  
2 #pragma omp parallel for \  
3   schedule(static,4) \  
4   num_threads(3)  
5   for (int j = 4; j < 20; j++) {  
6     b[i][j] = b[i][j-4] + x;  
7 #pragma omp barrier  
8   }  
9 }
```

Listing 11: OpenMP worksharing **for** construct with barrier (incorrect parallelization)

Incorrect parallelization! Barrier can not be placed inside a **parallel for** construct.

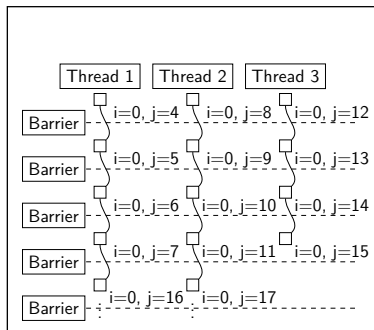


Figure 11: Possible OpenMP thread execution schedule for **parallel for** construct with **chunking**

Correct Parallelization Example:

```
1 #pragma omp parallel for \  
2     schedule(static,4) \  
3     num_threads(3) \  
4 for (int i = 0; i < 20; i++) { \  
5     for (int j = 4; j < 20; j++) { \  
6         b[i][j] = b[i][j-4] + x; \  
7     } \  
8 }
```

Listing 12: OpenMP worksharing **for** construct

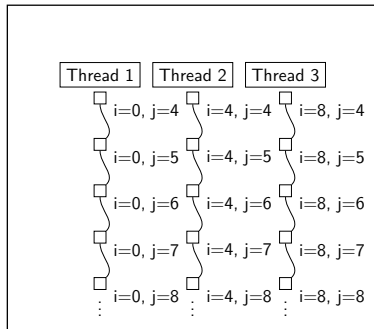


Figure 12: OpenMP thread execution schedule for **parallel for** construct with **chunking**

Correct Parallelization and Vectorization Example:

```
1 #pragma omp parallel for \  
2     schedule(static,4) \  
3     num_threads(3)  
4 for (int i = 0; i < 20; i++) {  
5 #pragma omp simd simdlen(4)  
6     for (int j = 4; j < 20; j++) {  
7         b[i][j] = b[i][j-4] + x;  
8     }  
9 }
```

Listing 13: OpenMP worksharing **for** construct

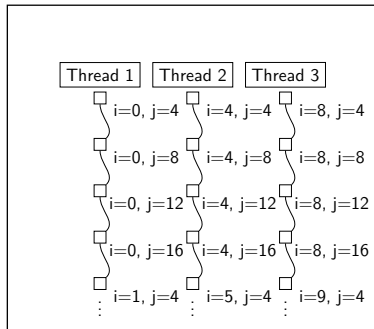


Figure 13: OpenMP thread execution schedule for **parallel for** construct with **chunking**

- 1 Introduction to OpenMP
- 2 What are Data Races?
- 3 Static data race detection in LLVM
- 4 Results
- 5 Hands-on with LLOV

Definition (Data Race)

An execution of a concurrent program is said to have a *data race* when two different threads access the same memory location,

- these accesses are not protected by a mutual exclusion mechanism
- the order of the two accesses is non-deterministic
- one of these accesses is a write

- Missing data sharing clauses

```
1 #pragma omp parallel for \  
2     private (temp,i,j)  
3     for (i = 0; i < len; i++)  
4         for (j = 0; j < len; j++){  
5             temp = u[i][j];  
6             sum = sum + temp * temp;  
7         }
```

Listing 14: DRB021: OpenMP Worksharing construct with data race

- Missing data sharing clauses
- Loop carried dependences

```
1  for (i=0;i<n;i++) {  
2  #pragma omp parallel for  
3      for (j=1;j<m;j++) {  
4      b[i][j] = b[i][j-1] ;  
5      }  
6  }
```

Listing 15: DRB038: Example with Loop Carried Dependence

- Missing data sharing clauses
- Loop carried dependences
- SIMD races

```
1 #pragma omp simd
2 for (int i=0; i<len-1; i++){
3     a[i+1] = a[i] + b[i];
4 }
```

Listing 16: DRB024: Example with SIMD data race

- Missing data sharing clauses
- Loop carried dependences
- SIMD races
- Synchronization issues

```
1 #pragma omp parallel \  
2     shared(b, error)  
3     {  
4 #pragma omp for nowait  
5     for(i = 0; i < len; i++)  
6         a[i] = b + a[i]*5;  
7 #pragma omp single  
8     error = a[9] + 1;  
9 }
```

Listing 17: DRB013: Example with data race due to improper synchronization

- Missing data sharing clauses
- Loop carried dependences
- SIMD races
- Synchronization issues
- Control flow dependent on number of threads

```
1 #pragma omp parallel
2   if (omp_get_thread_num() % 2 == 0) {
3     Flag = true;
4   }
```

Listing 18: Control flow dependent on number of threads

- 1 Introduction to OpenMP
- 2 What are Data Races?
- 3 Static data race detection in LLVM**
- 4 Results
- 5 Hands-on with LLOV

LLOV is a **language agnostic**, **static** OpenMP data race checker in the LLVM compiler framework.

LLOV is a **language agnostic**, **static** OpenMP data race checker in the LLVM compiler framework. LLOV

- is based on intermediate representation of LLVM (**LLVM-IR**)

LLOV is a **language agnostic**, **static** OpenMP data race checker in the LLVM compiler framework. LLOV

- is based on intermediate representation of LLVM (**LLVM-IR**)
- can handle FORTRAN as well as C/C++

LLOV is a **language agnostic**, **static** OpenMP data race checker in the LLVM compiler framework. LLOV

- is based on intermediate representation of LLVM (**LLVM-IR**)
- can handle FORTRAN as well as C/C++
- uses Polyhedral framework, **Polly**, of LLVM

LLOV is a **language agnostic**, **static** OpenMP data race checker in the LLVM compiler framework. LLOV

- is based on intermediate representation of LLVM (**LLVM-IR**)
- can handle FORTRAN as well as C/C++
- uses Polyhedral framework, **Polly**, of LLVM
- can conservatively state when a program is **data race free**

LLOV is a **language agnostic**, **static** OpenMP data race checker in the LLVM compiler framework. LLOV

- is based on intermediate representation of LLVM (**LLVM-IR**)
- can handle FORTRAN as well as C/C++
- uses Polyhedral framework, **Polly**, of LLVM
- can conservatively state when a program is **data race free**
- is capable of generating task graphs of OpenMP constructs

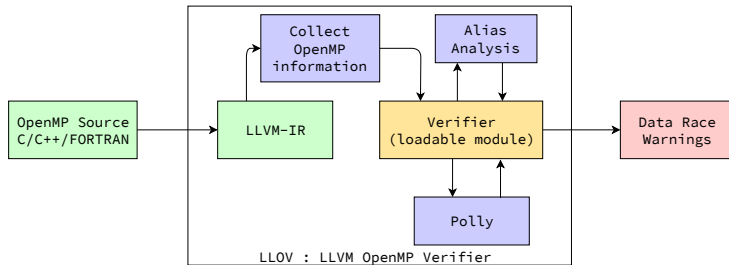


Figure 14: Flow Diagram of LLVM OpenMP Verifier (LLOV)

```
1 #pragma omp parallel shared(b, error)
2 {
3   #pragma omp for nowait
4   for(i = 0; i < len; i++)
5     a[i] = b + a[i]*5;
6   #pragma omp single
7     error = a[9] + 1;
8 }
```

```
1 #pragma omp parallel shared(b, error)
2 {
3 #pragma omp for nowait
4     for(i = 0; i < len; i++)
5         a[i] = b + a[i]*5;
6 #pragma omp single
7     error = a[9] + 1;
8 }
```

```
Directive:  OMP_Parallel
Variables:
  Private:   %.omp.ub = alloca i32, align 4
  Private:   %.omp.lb = alloca i32, align 4
  Shared:    i32* %i
  Shared:    i32* %len
  Firstprivate: i64 %vla
  Shared:    i32* %a
  Shared:    i32* %b
  Shared:    i32* %error
  Private:   %.omp.stride = alloca i32, align 4
  Private:   %.omp.is_last = alloca i32, align 4
Child Directives:
1: Directive:  OMP_Workshare_Loop
   Schedule type : Static Schedule (auto-chunked)
2: Directive:  OMP_Workshare_single
3: Directive:  OMP_Barrier
```

Listing 20: In-memory representation of a directive

```
1 #pragma omp parallel shared(b, error)
2 {
3 #pragma omp for nowait
4     for(i = 0; i < len; i++)
5         a[i] = b + a[i]*5;
6 #pragma omp single
7     error = a[9] + 1;
8 }
```

Directive: OMP_Parallel

Variables:

Private: %.omp.ub = alloca i32, align 4

Private: %.omp.lb = alloca i32, align 4

Shared: i32* %i

Shared: i32* %len

Firstprivate: i64 %vla

Shared: i32* %a

Shared: i32* %b

Shared: i32* %error

Private: %.omp.stride = alloca i32, align 4

Private: %.omp.is_last = alloca i32, align 4

Child Directives:

1: Directive: OMP_Workshare_Loop

Schedule type : Static Schedule (auto-chunked)

2: Directive: OMP_Workshare_single

3: Directive: OMP_Barrier

```
<Directive> ::= <Dtype> [ Sched ] { <Var> } { <Directive> }
<Dtype>      ::= parallel | for | simd
               | workshare | single
               | master | critical
<Var>        ::= <Vtype> val
<Vtype>      ::= private | firstprivate
               | shared | lastprivate
               | reduction | threadprivate
<Sched>      ::= [ <modifier> ] [ ordered ] <Stype> <chunk>
<modifier>   ::= monotonic | nonmonotonic
<Stype>      ::= static | dynamic | guided | auto | runtime
<chunk>      ::= positive-int-const
```

Listing 20: In-memory representation of a directive

```
1 for (i=0; i<m ; i++) {  
2 #pragma omp parallel for  
3     for (j=1; j<n ;j++) {  
4 SO:   b[i][j] = b[i][j-1];  
5     }  
6 }
```

Iteration Domain : $\mathbf{I} = \{\text{SO}(i,j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $\mathbf{S} = \{\text{SO}(i,j) \rightarrow (i,j)\} \cap_{dom} \mathbf{I}$

Access Map : $\mathbf{A} = \{\text{SO}(i,j) \rightarrow \mathbf{M}(i,j); \text{SO}(i,j) \rightarrow \mathbf{M}(i,j-1)\}$

Dependences : $\mathbf{D} = \{\text{SO}(i,j) \rightarrow (i,j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i=0; i<m ; i++) {  
2 #pragma omp parallel for  
3     for (j=1; j<n ;j++) {  
4 SO:   b[i][j] = b[i][j-1];  
5     }  
6 }
```

Iteration Domain : $\mathbf{I} = \{\text{SO}(i,j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $\mathbf{S} = \{\text{SO}(i,j) \rightarrow (i,j)\} \cap_{dom} \mathbf{I}$

Access Map : $\mathbf{A} = \{\text{SO}(i,j) \rightarrow \mathbf{M}(i,j); \text{SO}(i,j) \rightarrow \mathbf{M}(i,j-1)\}$

Dependences : $\mathbf{D} = \{\text{SO}(i,j) \rightarrow (i,j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i=0; i<m ; i++) {  
2 #pragma omp parallel for  
3     for (j=1; j<n ;j++) {  
4 SO:   b[i][j] = b[i][j-1];  
5     }  
6 }
```

Iteration Domain : $I = \{SO(i,j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $S = \{SO(i,j) \rightarrow (i,j)\} \cap_{dom} I$

Access Map : $A = \{SO(i,j) \rightarrow M(i,j); SO(i,j) \rightarrow M(i,j-1)\}$

Dependences : $D = \{SO(i,j) \rightarrow (i,j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$


```
1 for (i=0; i<m ; i++) {  
2 #pragma omp parallel for  
3   for (j=1; j<n ;j++) {  
4 SO:   b[i][j] = b[i][j-1];  
5   }  
6 }
```

Iteration Domain : $I = \{SO(i,j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $S = \{SO(i,j) \rightarrow (i,j)\} \cap_{dom} I$

Access Map : $A = \{SO(i,j) \rightarrow M(i,j); SO(i,j) \rightarrow M(i,j-1)\}$

Dependences : $D = \{SO(i,j) \rightarrow (i,j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i=0; i<m ; i++) {  
2 #pragma omp parallel for  
3   for (j=1; j<n ; j++) {  
4 SO:   b[i][j] = b[i][j-1];  
5   }  
6 }
```

Iteration Domain : $I = \{SO(i,j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $S = \{SO(i,j) \rightarrow (i,j)\} \cap_{dom} I$

Access Map : $A = \{SO(i,j) \rightarrow M(i,j); SO(i,j) \rightarrow M(i,j-1)\}$

Dependences : $D = \{SO(i,j) \rightarrow (i,j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i=0; i<m ; i++) {  
2 #pragma omp parallel for  
3   for (j=1; j<n ; j++) {  
4 SO:  b[i][j] = b[i][j-1];  
5   }  
6 }
```

Iteration Domain : $I = \{SO(i, j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $S = \{SO(i, j) \rightarrow (i, j)\} \cap_{dom} I$

Access Map : $A = \{SO(i, j) \rightarrow M(i, j); SO(i, j) \rightarrow M(i, j-1)\}$

Dependences : $D = \{SO(i, j) \rightarrow (i, j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i=0; i<m ; i++) {  
2 #pragma omp parallel for  
3     for (j=1; j<n ;j++) {  
4 SO:   b[i][j] = b[i][j-1];  
5     }  
6 }
```

Iteration Domain : $I = \{SO(i,j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $S = \{SO(i,j) \rightarrow (i,j)\} \cap_{dom} I$

Access Map : $A = \{SO(i,j) \rightarrow M(i,j); SO(i,j) \rightarrow M(i,j-1)\}$

Dependences : $D = \{SO(i,j) \rightarrow (i,j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i=0; i<m ; i++) {  
2 #pragma omp parallel for  
3   for (j=1; j<n ;j++) {  
4 SO:  b[i][j] = b[i][j-1];  
5   }  
6 }
```

Iteration Domain : $I = \{SO(i, j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $S = \{SO(i, j) \rightarrow (i, j)\} \cap_{dom} I$

Access Map : $A = \{SO(i, j) \rightarrow M(i, j); SO(i, j) \rightarrow M(i, j-1)\}$

Dependences : $D = \{SO(i, j) \rightarrow (i, j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i=0; i<m ; i++) {  
2 #pragma omp parallel for  
3     for (j=1; j<n ;j++) {  
4 SO:   b[i][j] = b[i][j-1];  
5     }  
6 }
```

Iteration Domain : $I = \{SO(i, j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $S = \{SO(i, j) \rightarrow (i, j)\} \cap_{dom} I$

Access Map : $A = \{SO(i, j) \rightarrow M(i, j); SO(i, j) \rightarrow M(i, j-1)\}$

Dependences : $D = \{SO(i, j) \rightarrow (i, j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1  for (i=0;i<10;i++) {  
2  #pragma omp parallel for  
3      for (j=1;j<10;j++) {  
4          b[i][j]=b[i][j-1];  
5      }  
6  }
```

Listing 23: Example with Loop Carried
Dependence

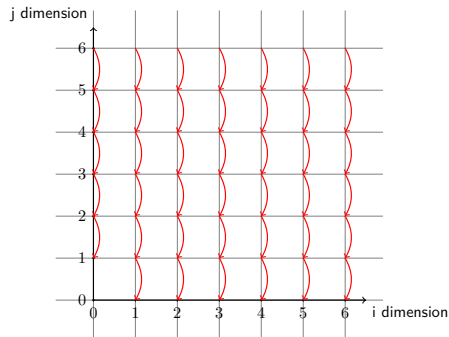


Figure 15: Dependence Polyhedra

```
1  for (i=0;i<10;i++) {  
2  #pragma omp parallel for  
3      for (j=1;j<10;j++) {  
4          b[i][j]=b[i][j-1];  
5      }  
6  }
```

Listing 24: Example with Loop Carried
Dependence

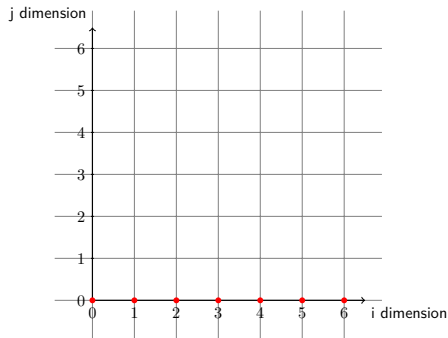


Figure 16: Projection of the Dependence Polyhedra on
i-dimension

Zero magnitude of the projections on a dimension signifies that the dimension is **parallel**.


```
1  for (i=0;i<10;i++) {  
2  #pragma omp parallel for  
3      for (j=1;j<10;j++) {  
4          b[i][j]=b[i][j-1];  
5      }  
6  }
```

Listing 25: Example with Loop Carried Dependence

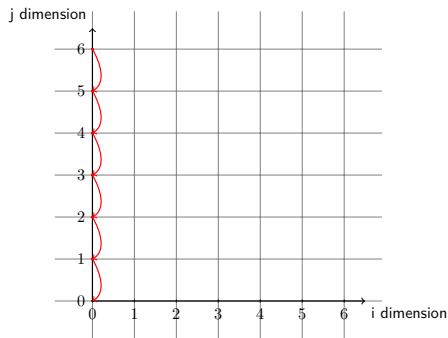


Figure 17: Projection of the Dependence Polyhedra on j-dimension

Non-zero magnitude of the projections on a dimension signifies that the dimension is **not** parallel.

- 1 Introduction to OpenMP
- 2 What are Data Races?
- 3 Static data race detection in LLVM
- 4 Results**
- 5 Hands-on with LLOV

Benchmarks:

- DataRaceBench C/C++ v1.2 [Liao et al., 2018a, Liao et al., 2018b]
- OmpSCR v2.0 [Dorta et al., 2004, Dorta et al., 2005]
- DataRaceBench FORTRAN [Kukreja et al., 2019]

Benchmarks:

- DataRaceBench C/C++ v1.2 [Liao et al., 2018a, Liao et al., 2018b]
- OmpSCR v2.0 [Dorta et al., 2004, Dorta et al., 2005]
- DataRaceBench FORTRAN [Kukreja et al., 2019]

System Specifications:

System: Two Intel Xeon E5-2697 v4 @ 2.30GHz processors

OS: 64 bit Ubuntu 18.04.2 LTS server

Kernel: Linux kernel version 4.15.0-48-generic

Threads: 72 (2 x 36) hardware threads

Memory: 128GB

OpenMP library: LLVM OpenMP runtime v5.0.1 (libomp5)

Table 3: Race detection tools with the version numbers used for comparison

Tools	Source	Version / Commit
HELGRIND [Valgrind-project, 2007b]	Valgrind	3.13.0
VALGRIND DRD [Valgrind-project, 2007a]	Valgrind	3.13.0
TSAN-LLVM [Serebryany and Iskhodzhanov, 2009]	LLVM	6.0.1
ARCHER [Atzeni et al., 2016]	git master branch	fc17353
SWORD [Atzeni et al., 2018]	git master branch	7a08f3c

Results: DataRaceBench v1.2 comparison



Table 4: Maximum number of Races reported by different tools in DataRaceBench 1.2

Tools	Race: Yes		Race: No		Coverage/116
	TP	FN	TN	FP	
HELGRIND	56	3	2	55	116
VALGRIND DRD	56	3	26	31	116
TSAN-LLVM	57	2	2	55	116
ARCHER	56	3	2	55	116
SWORD	47	4	24	4	79
LLOV	48	2	36	5	91

Results: DataRaceBench v1.2 comparison



Table 4: Maximum number of Races reported by different tools in DataRaceBench 1.2

Tools	Race: Yes		Race: No		Coverage/116
	TP	FN	TN	FP	
HELGRIND	56	3	2	55	116
VALGRIND DRD	56	3	26	31	116
TSAN-LLVM	57	2	2	55	116
ARCHER	56	3	2	55	116
SWORD	47	4	24	4	79
LLOV	48	2	36	5	91

Table 5: Maximum number of Races reported by different tools in common 61 kernels of DataRaceBench 1.2

Tools	Race: Yes		Race: No		Coverage/61
	TP	FN	TN	FP	
HELGRIND	42	1	2	16	61
VALGRIND DRD	42	1	12	6	61
TSAN-LLVM	42	1	2	16	61
ARCHER	42	1	2	16	61
SWORD	42	1	17	1	61
LLOV	42	1	16	2	61

Table 6: Performance of the tools on DataRaceBench 1.2

Tools	Precision	Recall	Accuracy	F1 Score	Diagnostic odds ratio
HELGRIND	0.50	0.95	0.50	0.66	0.68
VALGRIND DRD	0.64	0.95	0.71	0.77	15.66
TSAN-LLVM	0.51	0.97	0.51	0.67	1.04
ARCHER	0.50	0.95	0.50	0.66	0.68
SWORD	0.92	0.92	0.90	0.92	70.50
LLOV	0.91	0.96	0.92	0.93	172.80

Table 6: Performance of the tools on DataRaceBench 1.2

Tools	Precision	Recall	Accuracy	F1 Score	Diagnostic odds ratio
HELGRIND	0.50	0.95	0.50	0.66	0.68
VALGRIND DRD	0.64	0.95	0.71	0.77	15.66
TSAN-LLVM	0.51	0.97	0.51	0.67	1.04
ARCHER	0.50	0.95	0.50	0.66	0.68
SWORD	0.92	0.92	0.90	0.92	70.50
LLOV	0.91	0.96	0.92	0.93	172.80

Table 7: Performance of the tools on common 61 kernels of DataRaceBench 1.2

Tools	Precision	Recall	Accuracy	F1 Score	Diagnostic odds ratio
HELGRIND	0.72	0.98	0.72	0.83	5.25
VALGRIND DRD	0.88	0.98	0.89	0.92	84.00
TSAN-LLVM	0.72	0.98	0.72	0.83	5.25
ARCHER	0.72	0.98	0.72	0.83	5.25
SWORD	0.98	0.98	0.97	0.98	714.00
LLOV	0.95	0.98	0.95	0.97	336.00

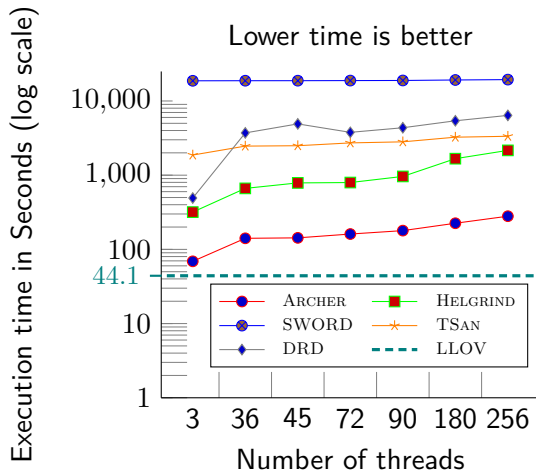


Figure 18: DataRaceBench v1.2 total execution time by different tools on logarithmic scale

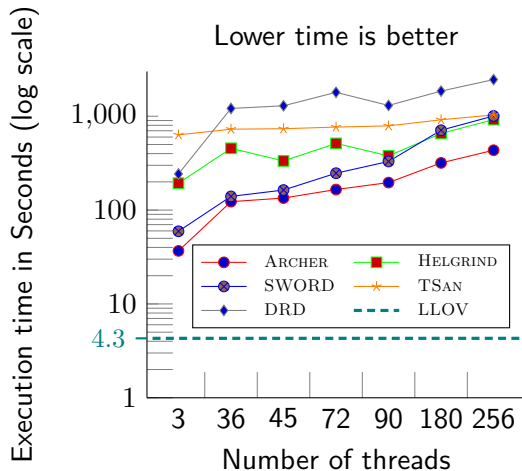


Figure 19: DataRaceBench v1.2 total time taken by different tools for common 61 kernels on logarithmic scale

Table 8: Comparison of different tools on OmpSCR v2.0

Tools	Race: Yes		Race: No		Coverage/14
	TP	FN	TN	FP	
HELGRIND	8	0	0	9	14
VALGRIND DRD	8	0	2	5	14
TSAN-LLVM	7	1	2	6	14
ARCHER	7	1	2	4	14
SWORD	3	4	3	0	10
LLOV	4	1	2	5	10

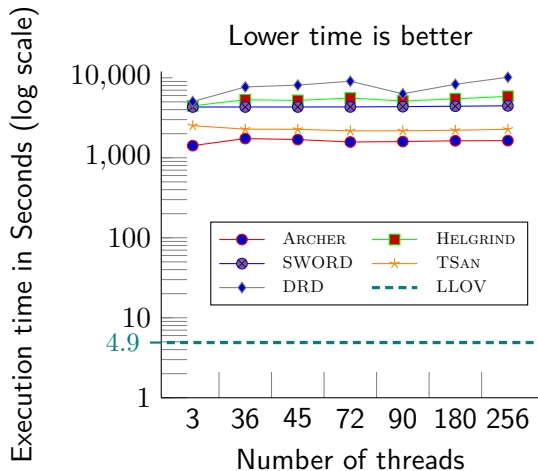


Figure 20: OmpSCR v2.0 total execution time by different tools on logarithmic scale

An implementation of DataRaceBench C/C++ v1.2 [Liao et al., 2018b] in FORTRAN 95.

- Converted 92 (out of 116) C/C++ kernels to FORTRAN
- Demonstrate that LLOV is language agnostic
- Already open-sourced this benchmark [Kukreja et al., 2019]

Table 9: Maximum number of Races reported by different tools in DataRaceBench FORTRAN

Tools	Race: Yes		Race: No		Coverage/92
	TP	FN	TN	FP	
HELGRIND	46	6	4	36	92
VALGRIND DRD	45	7	21	19	92
LLOV	36	7	19	5	67

LLOV: A Fast Static Data-Race Checker for OpenMP Programs

LLOV is freely available for download.

Link: <https://github.com/utpalbora/llov>

Blog: <https://compilers.cse.iith.ac.in/projects/llov/>



utpalbora.com

Open source links:

- DataRaceBench FORTRAN: https://github.com/IITH-Compilers/drb_fortran
- LLOV source: Please drop me an email at **cs14mtech11017@iith.ac.in**

Contributions Welcome!!

We welcome your contributions in any form. **Thank You!**

- 1 Introduction to OpenMP
- 2 What are Data Races?
- 3 Static data race detection in LLVM
- 4 Results
- 5 Hands-on with LLOV**



Atzeni, S., Gopalakrishnan, G., Rakamaric, Z., Ahn, D. H., Laguna, I., Schulz, M., Lee, G. L., Protze, J., and Müller, M. S. (2016).

Archer: effectively spotting data races in large openmp applications.

In *Parallel and Distributed Processing Symposium, 2016 IEEE International*, pages 53–62, Chicago, IL, USA. IEEE, IEEE.



Atzeni, S., Gopalakrishnan, G., Rakamaric, Z., Laguna, I., Lee, G. L., and Ahn, D. H. (2018).

Sword: A bounded memory-overhead detector of openmp data races in production runs.

In *2018 IEEE International Parallel and Distributed Processing Symposium (IPDPS)*, pages 845–854, Vancouver, BC, Canada. IEEE, IEEE.



Dorta, A. J., Rodriguez, C., and de Sande, F. (2004).

OpenMP Source Code Repository.

<https://sourceforge.net/projects/ompscr/files/OmpSCR/>.

[Online; accessed 19-May-2019].



Dorta, A. J., Rodriguez, C., and de Sande, F. (2005).

The openmp source code repository.

In *13th Euromicro Conference on Parallel, Distributed and Network-Based Processing*, pages 244–250, Washington, DC, USA. IEEE Computer Society.



Kukreja, P., Shukla, H., and Bora, U. (2019).
DataRaceBench FORTRAN.
https://github.com/IITH-Compilers/drb_fortran.
[Online; accessed 19-October-2019].



Liao, C., Lin, P.-H., Asplund, J., Schordan, M., and Karlin, I. (2018a).
DataRaceBench v1.2.0.
<https://github.com/LLNL/dataracebench>.
[Online; accessed 19-May-2019].



Liao, C., Lin, P.-H., Schordan, M., and Karlin, I. (2018b).
A semantics-driven approach to improving dataracebench's openmp standard coverage.
In *Evolving OpenMP for Evolving Architectures*, pages 189–202, Cham. Springer International Publishing.



Mayer, F., Knaust, M., and Philippsen, M. (2019).
Openmp on fpgas—a survey.
In Fan, X., de Supinski, B. R., Sinnen, O., and Giacaman, N., editors, *OpenMP: Conquering the Full Hardware Spectrum*, pages 94–108, Cham. Springer International Publishing.



Serebryany, K. and Iskhodzhanov, T. (2009).

Threadsanitizer: Data race detection in practice.

In *Proceedings of the Workshop on Binary Instrumentation and Applications*, WBIA '09, pages 62–71, New York, NY, USA. ACM.



Valgrind-project (2007a).

DRD: a thread error detector.

<http://valgrind.org/docs/manual/drd-manual.html>.

[Online; accessed 08-May-2019].



Valgrind-project (2007b).

Helgrind: a thread error detector.

<http://valgrind.org/docs/manual/hg-manual.html>.

[Online; accessed 08-May-2019].

Extra Slides!

Algorithm 1: Race Detection Algorithm

Input: Loop L

Output: True/False

```
1 Function isRaceFree( $L$ ):  
2    $SCoP$  = ConstructSCoP( $L$ ) ;  
3    $RDG$  = ComputeDependences( $SCoP$ ) ;  
4    $depth$  = GetLoopDepth( $L$ ) ;  
5   if isParallel( $RDG$ ,  $depth$ ) then  
6     // Program is race free.  
7     return True ;  
8   else  
9     // Data Race detected.  
10    return False ;  
11  return result  
12 End Function
```

Algorithm 2: Algorithm to check parallelism

Input: RDG , Loop-depth dim

Output: True/False

```
1 Function isParallel( $RDG$ ,  $dim$ ):  
2   if  $RDG$  is Empty then  
3     return True ;  
4   else  
5     Flag = True;  
6     while Dependence  $D$  in  $RDG$  do  
7        $D'$  = Project Out all dimensions except  $dim$  from  $D$  ;  
8       if  $D'$  is Empty then  
9         continue ;  
10      else  
11        Flag = False ;  
12        break ;  
13      return Flag ;  
14 End Function
```

Table 10: Comparison of OpenMP pragma handling by OpenMP aware tools. (Y for Yes, N for No)

OpenMP Pragma	LLOV	OMPVERIFY	POLYOMP	DRACO	SWORD	ARCHER	ROMP
#pragma omp parallel	Y	Y	Y	Y	Y	Y	Y
#pragma omp for	Y	Y	Y	Y	Y	Y	Y
#pragma omp parallel for	Y	Y	Y	Y	Y	Y	Y
#pragma omp critical	N	N	N	N	Y	Y	Y
#pragma omp atomic	N	N	N	N	Y	Y	Y
#pragma omp master	N	N	Y	N	Y	Y	Y
#pragma omp single	N	N	Y	N	Y	Y	Y
#pragma omp simd	Y	N	N	Y	N	N	N
#pragma omp parallel for simd	Y	N	N	Y	N	N	N
#pragma omp parallel sections	N	N	N	N	Y	Y	Y
#pragma omp sections	N	N	N	N	Y	Y	Y
#pragma omp threadprivate	Y	N	N	N	N	Y	Y
#pragma omp ordered	Y	N	N	N	N	Y	Y
#pragma omp distribute	Y	N	N	N	N	Y	Y
#pragma omp task	N	N	N	N	N	Y	Y
#pragma omp taskgroup	N	N	N	N	N	Y	Y
#pragma omp taskloop	N	N	N	N	N	Y	Y
#pragma omp taskwait	N	N	N	N	N	Y	Y
#pragma omp barrier	N	N	Y	N	Y	Y	Y
#pragma omp teams	N	N	N	N	N	N	N
#pragma omp target	N	N	N	N	N	N	N
#pragma omp target map	N	N	N	N	N	N	N

Table 11: Race detection literature (minimalistic)

Published	Title	Year
TOCS	Eraser: A dynamic data race detector for multithreaded programs	1997
SOSP	RacerX: Effective, Static Detection of Race Conditions and Deadlocks	2003
POPL	Conditional must not aliasing for static race detection	2007
PPoPP	May-happen-in-parallel analysis of X10 program	2007
ESEC-FSE	RELAY: static race detection on millions of lines of code	2007
WBIAS	ThreadSanitizer – data race detection in practice	2009
IWOMP	ompVerify : Polyhedral Analysis for the OpenMP Programmer	2011
TOPLAS	LOCKSMITH: Practical static race detection for C	2011
LCPC	An extended polyhedral model for SPMD programs and its use in static data race detection	2016
IPDPS	ARCHER: effectively spotting data races in large OpenMP applications	2016
IPDPS	SWORD: A Bounded Memory-Overhead Detector of OpenMP Data Races in Production Runs	2018
Correctness	Using Polyhedral Analysis to Verify OpenMP Applications are Data Race Free	2018
SC	Dynamic Data Race Detection for OpenMP Programs	2018
OOPSLA	RacerD: Compositional Static Race Detection	2018
IWOMP	OMPSan: Static Verification of OpenMP's Data Mapping Constructs	2019
SC	OMPRacer: A Scalable and Precise Static Race Detector for OpenMP Programs	2020

- **True Positive (TP):** If the evaluation tool correctly detects a data race present in the kernel it is a True Positive test result. A higher number of true positives represents a better tool.
- **True Negative (TN):** If the benchmark does not contain a race and the tool declares it as race-free, then it is a true negative case. A higher number of true negatives represents a better tool.
- **False Positives (FP):** If the benchmark does not contain any race, but the tool reports a race condition, it is a false positive. False Positives should be as low as possible.
- **False Negatives (FN):** False Negative test result is obtained when the tool fails to detect a known race in the benchmark. These are the cases that are missed by the tool. A lower number of false negatives are desirable.

- **Precision** : Precision is the measure of closeness of the outcomes of prediction. Thus, a higher value of precision represents that the tool will more often than not identify a race condition when it exists.

$$Precision = \frac{TP}{TP + FP}$$

- **Recall** : Recall gives the total number of cases detected out of the maximum data races present. A higher recall value means that there are less chances that a data race is missed by the tool. It is also called true positive rate (TPR).

$$Recall = \frac{TP}{TP + FN}$$

- **Accuracy** : Accuracy gives the chances of correct reports out of all the reports, as the name suggests. A higher value of accuracy is always desired and gives overall measure of the efficacy of the tool.

$$Accuracy = \frac{TP + TN}{TP + FP + TN + FN}$$

- **F1 Score** : The harmonic mean of precision and recall is called the F1 score. An F1 score of 1 can be achieved in the best case when both precision and recall are perfect. The worst case F1 score is 0 when either precision or recall is 0.

$$F1\ Score = 2 * \frac{Precision * Recall}{Precision + Recall}$$

- **Diagnostic odds ratio (DOR)** : It is the ratio of the positive likelihood ratio (LR+) to the negative likelihood ratio (LR-).

$$DOR = \frac{LR+}{LR-} \text{ where,}$$

$$\text{Positive Likelihood Ratio (LR+)} = \frac{TPR}{FPR},$$

$$\text{Negative Likelihood Ratio (LR-)} = \frac{FNR}{TNR},$$

$$\text{True Positive Rate (TPR)} = \frac{TP}{TP + FN},$$

$$\text{False Positive Rate (FPR)} = \frac{FP}{FP + TN},$$

$$\text{False Negative Rate (FNR)} = \frac{FN}{FN + TP} \text{ and}$$

$$\text{True Negative Rate (TNR)} = \frac{TN}{TN + FP}$$

DOR is the measure of the ratio of the odds of race detection being positive given that the test case has a data race, to the odds of race detection being positive given the test case does not have a race.