

Techniques for Static Detection of Data-Races in OpenMP Programs



Utpal Bora

PhD Scholar
Computer Science and Engineering
IIT Hyderabad, India

April 25, 2022

Doctoral Advisory Committee:

- Prof. Ramakrishna Upadrasta (Guide)
- Prof. Saurabh Joshi (Co-Guide)
- Prof. M.V. Panduranga Rao
- Prof. Sparsh Mittal
- Prof. Subrahmanyam Kalyanasundaram

Mentors:

- Dr. Venugopal Raghavan
- Dr. Johannes Doerfert
- Dr. Michael Kruse
- Prof. Tobias Grosser
- Prof. Krishna Nandivada

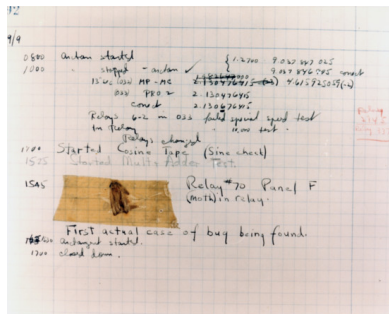
Collaborators:

- Prof. Sanjay Rajopadhye
- Santanu Das
- Pankaj Kukreja
- Himanshu Shukla
- Shraiys Vaishay

Problem Statement

“Bug” or “Glitch” or “Anomaly” or “Defect” is the incorrectness in software behaviour.

“Bug” word is popularized by Grace Hopper when she traced a problem in the Harvard Mark II system to a Moth trapped in the system.



Source: Wikipedia



Major Incidents due to Software Bugs



Bugs can result in “Fault”/“Error” or “Failure”/“Crash”.

Major Incidents due to Software Bugs

Bugs can result in “Fault”/“Error” or “Failure”/“Crash”.



Source: Wikipedia

- Ariane Flight V88 Explosion
- Mariner 1 Launch Failure
- Therac-25 Radiation Therapy Accident (lives lost)
- Northeast Blackout of 2003
- Heathrow Terminal 5 Opening Blunder (500 flights cancelled)

Problem Statement




Provide a solution to the problem of statically detecting **data-races** in shared-memory parallel programs.

Problem Statement

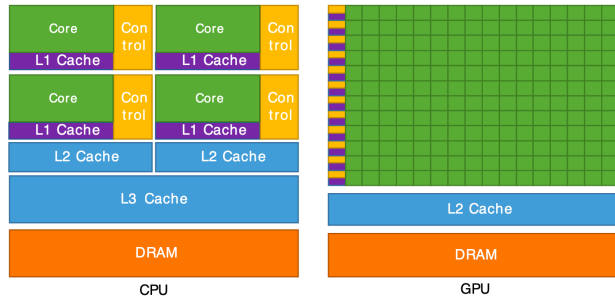


Provide a solution to the problem of statically detecting **data-races** in shared-memory parallel programs.



Parallel Programming Models

- Shared Memory
- Message Passing
- Partitioned Global Address Space (PGAS)



Source: https://cww.cac.cornell.edu/GPUarch/gpu_characteristics

Shared-Memory Programming Languages: Cilk, Cilk Plus, C++11, C#, CUDA, Java, **OpenMP**, Threading Building Blocks, etc.

Open Multi-Processing (OpenMP) is a framework for C, C++ and FORTRAN programs to enable **shared-memory** parallel programming on a multi-threaded system.

OpenMP **Supports:**


- *Single Program Multiple Data* (**SMPD**) programming model
- *Single Instruction Multiple Data* (**SIMD**) programming model
- **Divide-and-conquer** paradigm with tasks
- **Data environment** for shared-memory consistency
- **Mutual exclusion** and **atomicity**
- Implicit and explicit **synchronization**

OpenMP is **portable** across architectures from different vendors. (e.g. CPU, GPU, and FPGA [Mayer et al., 2019])

OpenMP Execution Model: parallel construct

```
1 #pragma omp parallel \  
2     num_threads(3)  
3 {  
4     for (int i = 0; i < 15; i++) {  
5         printf("Iteration Id %d by \  
6             Thread Id: %d\n", i,  
7             omp_get_thread_num());  
8     }  
9 }
```

OpenMP parallel construct




OpenMP thread execution model for parallel construct

OpenMP Execution Model: parallel for construct

```
1 #pragma omp parallel for \  
2     schedule(static) \  
3     num_threads(3)  
4 for (int i = 0; i < 15; i++) {  
5     printf("Iteration Id %d by \  
6         Thread Id: %d\n", i,  
7         omp_get_thread_num());  
8 }
```

OpenMP worksharing for construct



OpenMP thread execution model for parallel for construct

What is a Data-Race?



```
1 #pragma omp parallel \  
2     shared(b, error)  
3 {  
4 #pragma omp for nowait  
5     for(i = 0; i < len; i++)  
6         a[i] = b + a[i] * 5;  
7 #pragma omp single  
8     error = a[9] + 1;  
9 }
```

DRB013: Example with **data-race** due to improper synchronization

Definition (Data-Race)

An execution of a concurrent program is said to have a **data-race** when

- two threads access the **same memory location**,
- accesses are **not** protected by a **mutual exclusion** mechanism (e.g., locks),
- the order of the two accesses is **non-deterministic**,
- one of these accesses is a **write**.

- Missing data sharing clauses

```
1 #pragma omp parallel for \  
2     private (temp,i,j)  
3     for (i = 0; i < len; i++)  
4         for (j = 0; j < len; j++){  
5             temp = u[i][j];  
6             sum = sum + temp * temp;  
7         }
```

DRB021: OpenMP Worksharing construct with `data-race`

- Missing data sharing clauses
- Loop carried dependences

```
1  for (i = 0; i < n; i++) {  
2  #pragma omp parallel for  
3  for (j = 1; j < m; j++) {  
4      b[i][j] = b[i][j-1];  
5  }  
6  }
```

DRB038: Example with Loop Carried Dependence

Types of data-races in OpenMP programs



- Missing data sharing clauses
- Loop carried dependences
- SIMD races

```
1 #pragma omp simd
2 for (int i = 0; i < len-1; i++){
3     a[i+1] = a[i] + b[i];
4 }
```

DRB024: Example with SIMD data-race

- Missing data sharing clauses
- Loop carried dependences
- SIMD races
- Synchronization issues

```
1 #pragma omp parallel \  
2     shared(b, error)  
3 {  
4 #pragma omp for nowait  
5     for(i = 0; i < len; i++)  
6         a[i] = b + a[i] * 5;  
7 #pragma omp single  
8     error = a[9] + 1;  
9 }
```

DRB013: Example with data-race due to improper synchronization

Types of data-races in OpenMP programs



- Missing data sharing clauses
- Loop carried dependences
- SIMD races
- Synchronization issues
- Control flow dependent on number of threads

```
1 #pragma omp parallel \  
2     shared(Flag)  
3 {  
4     if (omp_get_thread_num() % 2 == 0) {  
5         Flag = true;  
6     }  
7     //  
8 }
```

Control flow dependent on number of threads

- Using Polyhedral Dependence Analysis
- Using May-Happen-in-Parallel (MHP) Analysis

Race Detection using Polyhedral Dependence Analysis

```
1 for (i = 0; i < m; i++) {  
2   #pragma omp parallel for  
3     for (j = 1; j < n; j++) {  
4   S0:   b[i][j] = b[i][j-1];  
5     }  
6   }
```

Iteration Domain : $I = \{S0(i, j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $S = \{S0(i, j) \rightarrow (i, j)\} \cap_{dom} I$

Access Map : $A = \{S0(i, j) \rightarrow M(i, j); S0(i, j) \rightarrow M(i, j-1)\}$

Dependences : $D = \{S0(i, j) \rightarrow (i, j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i = 0; i < m; i++) {  
2   #pragma omp parallel for  
3   for (j = 1; j < n; j++) {  
4     S0: b[i][j] = b[i][j-1];  
5   }  
6 }
```

Iteration Domain : $I = \{S0(i, j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $S = \{S0(i, j) \rightarrow (i, j)\} \cap_{dom} I$

Access Map : $A = \{S0(i, j) \rightarrow M(i, j); S0(i, j) \rightarrow M(i, j-1)\}$

Dependences : $D = \{S0(i, j) \rightarrow (i, j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i = 0; i < m; i++) {  
2   #pragma omp parallel for  
3   for (j = 1; j < n; j++) {  
4     S0: b[i][j] = b[i][j-1];  
5   }  
6 }
```

Iteration Domain : $I = \{S0(i, j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $S = \{S0(i, j) \rightarrow (i, j)\} \cap_{dom} I$

Access Map : $A = \{S0(i, j) \rightarrow M(i, j); S0(i, j) \rightarrow M(i, j-1)\}$

Dependences : $D = \{S0(i, j) \rightarrow (i, j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i = 0; i < m; i++) {  
2   #pragma omp parallel for  
3   for (j = 1; j < n; j++) {  
4     S0: b[i][j] = b[i][j-1];  
5   }  
6 }
```

Iteration Domain : $I = \{S0(i, j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $S = \{S0(i, j) \rightarrow (i, j)\} \cap_{dom} I$

Access Map : $A = \{S0(i, j) \rightarrow M(i, j); S0(i, j) \rightarrow M(i, j-1)\}$

Dependences : $D = \{S0(i, j) \rightarrow (i, j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$


```
1 for (i = 0; i < m; i++) {  
2   #pragma omp parallel for  
3   for (j = 1; j < n; j++) {  
4     S0: b[i][j] = b[i][j-1];  
5   }  
6 }
```

Iteration Domain : $I = \{S0(i, j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $S = \{S0(i, j) \rightarrow (i, j)\} \cap_{dom} I$

Access Map : $A = \{S0(i, j) \rightarrow M(i, j); S0(i, j) \rightarrow M(i, j-1)\}$

Dependences : $D = \{S0(i, j) \rightarrow (i, j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i = 0; i < m; i++) {  
2   #pragma omp parallel for  
3   for (j = 1; j < n; j++) {  
4   S0: b[i][j] = b[i][j-1];  
5   }  
6 }
```

Iteration Domain : $I = \{S0(i, j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $S = \{S0(i, j) \rightarrow (i, j)\} \cap_{dom} I$

Access Map : $A = \{S0(i, j) \rightarrow M(i, j); S0(i, j) \rightarrow M(i, j-1)\}$

Dependences : $D = \{S0(i, j) \rightarrow (i, j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i = 0; i < m; i++) {  
2   #pragma omp parallel for  
3   for (j = 1; j < n; j++) {  
4     S0: b[i][j] = b[i][j-1];  
5   }  
6 }
```

Iteration Domain : $I = \{S0(i, j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $S = \{S0(i, j) \rightarrow (i, j)\} \cap_{dom} I$

Access Map : $A = \{S0(i, j) \rightarrow M(i, j); S0(i, j) \rightarrow M(i, j-1)\}$

Dependences : $D = \{S0(i, j) \rightarrow (i, j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i = 0; i < m; i++) {  
2   #pragma omp parallel for  
3   for (j = 1; j < n; j++) {  
4   S0: b[i][j] = b[i][j-1];  
5   }  
6 }
```

Iteration Domain : $I = \{S0(i, j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $S = \{S0(i, j) \rightarrow (i, j)\} \cap_{dom} I$

Access Map : $A = \{S0(i, j) \rightarrow M(i, j); S0(i, j) \rightarrow M(i, j-1)\}$

Dependences : $D = \{S0(i, j) \rightarrow (i, j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

```
1 for (i = 0; i < m; i++) {  
2   #pragma omp parallel for  
3     for (j = 1; j < n; j++) {  
4   S0:   b[i][j] = b[i][j-1];  
5     }  
6   }
```

Iteration Domain : $I = \{S0(i, j) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Schedule : $S = \{S0(i, j) \rightarrow (i, j)\} \cap_{dom} I$


Access Map : $A = \{S0(i, j) \rightarrow M(i, j); S0(i, j) \rightarrow M(i, j-1)\}$

Dependences : $D = \{S0(i, j) \rightarrow (i, j-1) : 0 \leq i \leq m-1 \wedge 1 \leq j \leq n-1\}$

Technique 1: Race Detection using Polyhedral Dependences

```
1  for (int i = 0; i < m; i++) {  
2  #pragma omp parallel for  
3      for (int j = 1; j < n; j++) {  
4          b[i][j] = b[i][j-1];  
5      }  
6  }
```

Example with Loop Carried Dependence




Dependence Polyhedra

Technique 1: Race Detection using Polyhedral Dependences

```
1  for (int i = 0; i < m; i++) {  
2  #pragma omp parallel for  
3      for (int j = 1; j < n; j++) {  
4          b[i][j] = b[i][j-1];  
5      }  
6  }
```

Example with Loop Carried Dependence




Projection of the Dependence Polyhedra on i-dimension

Zero magnitude of the projections on a dimension signifies that the dimension is parallel.

Technique 1: Race Detection using Polyhedral Dependences

```
1  for (int i = 0; i < m; i++) {  
2  #pragma omp parallel for  
3      for (int j = 1; j < n; j++) {  
4          b[i][j] = b[i][j-1];  
5      }  
6  }
```

Example with Loop Carried Dependence



Projection of the Dependence Polyhedra on j-dimension

Non-zero magnitude of the projections on a dimension signifies that the dimension is **not** parallel.

- Limited scope of the Polyhedral framework (less coverage).

```
1 void f1(int* q) {  
2     *q += 1;  
3     // None-affine code  
4 }  
5  
6 int main() {  
7     int i = 0;  
8     #pragma omp parallel  
9     { f1(&i); }  
10 }
```

Non-affine program

Overcoming Limitations of the Polyhedral Framework



- Limited scope of the Polyhedral framework (less coverage).
- Cannot handle inter SCoP (Static Control Parts) dependencies.

```
1 void f1(int* q) {  
2     *q += 1;  
3     // None-affine code  
4 }  
5  
6 int main() {  
7     int i = 0;  
8     #pragma omp parallel  
9     { f1(&i); }  
10 }
```

Non-affine program

```
1 #pragma omp parallel \  
2     shared(b, error)  
3 {  
4     #pragma omp for nowait  
5     for(i = 0; i < len; i++)  
6         a[i] = b + a[i]*5;  
7     #pragma omp single  
8     error = a[9] + 1;  
9 }
```

DRB013: Example with data-race due to improper synchronization

Overcoming Limitations of the Polyhedral Framework



- Limited scope of the Polyhedral framework (less coverage). Use [Dependence Testing](#)
- Cannot handle inter SCoP (Static Control Parts) dependencies. Use [MHP Analysis](#)

```
1 void f1(int* q) {  
2     *q += 1;  
3     // None-affine code  
4 }  
5  
6 int main() {  
7     int i = 0;  
8     #pragma omp parallel  
9     { f1(&i); }  
10 }
```

Non-affine program

```
1 #pragma omp parallel \  
2     shared(b, error)  
3 {  
4     #pragma omp for nowait  
5     for(i = 0; i < len; i++)  
6         a[i] = b + a[i]*5;  
7     #pragma omp single  
8     error = a[9] + 1;  
9 }
```

DRB013: Example with [data-race](#) due to improper synchronization

Race Detection using May-Happen-in-Parallel Analysis

- Perform **Phase Interval Analysis (PIA)** in the Monotone framework [Kildall, 1973]
 - Construct a **TASKGRAPH** for the OpenMP constructs
 - Compute Phase Intervals ($[lb, ub]$) for each basic block
- Perform **May-Happen-in-Parallel (MHP)** analysis using the Phase Intervals
 - Compute MHP sets for each basic block
- Use MHP sets to detect **data-races** statically

Modeling OpenMP single construct as TASKGRAPH




```
1 #pragma omp parallel
2 {
3     // S1
4 #pragma omp single
5     {
6         // S2
7     } // implicit barrier
8     // S3
9 #pragma omp single
10    {
11        // S4
12    } // implicit barrier
13    // S5
14 }
```

Multiple OpenMP single constructs

Modeling OpenMP single construct as TASKGRAPH

```
1 #pragma omp parallel
2 {
3     // S1
4 #pragma omp single
5 {
6     // S2
7 } // implicit barrier
8 // S3
9 #pragma omp single
10 {
11     // S4
12 } // implicit barrier
13 // S5
14 }
```

Multiple OpenMP single constructs



TASKGRAPH for the OpenMP single construct annotated with phase intervals.

Modeling OpenMP master construct as TASKGRAPH



```
1 #pragma omp parallel
2 {
3     // S1
4 #pragma omp master
5 {
6     // S2
7 }
8 // S3
9 #pragma omp master
10 {
11     // S4
12 }
13 // S5
14 }
```


Multiple OpenMP master constructs

Modeling OpenMP master construct as TASKGRAPH



```
1 #pragma omp parallel
2 {
3     // S1
4 #pragma omp master
5 {
6     // S2
7 }
8 // S3
9 #pragma omp master
10 {
11     // S4
12 }
13 // S5
14 }
```

Multiple OpenMP master constructs



TASKGRAPH for the OpenMP master construct annotated with phase intervals.

Definition (Reduced TASKGRAPH)

A reduced TASKGRAPH $G = \langle V, E, R, T \rangle$ consists of:

- the vertex set V of implicit and explicit tasks.
- the edge set $E \subseteq V \times V$ contains directed edges.
- a special vertex R for the root of the graph.
- a special sentinel vertex T that represents the termination of all the tasks.

Definition (Phase)

A phase, represented by a positive integer $p \in \mathbb{N}^+ \cup \infty$, is the execution number of an instance of the nodes in the `TASKGRAPH`.


A program can execute in phases due to:

- the presence of barriers
- upon entry to a `parallel` region
- upon exit from a `parallel` region.

Formalizing Phase Interval Analysis (PIA)

Lattice : $(I, \sqsubseteq, \sqcup, \sqcap)$ is a lattice over an interval domain I .

Top element: $\top = [0, \infty]$. Bot element: \perp .




Interval Lattice for PIA

Formalizing Phase Interval Analysis (PIA)



Lattice : $(I, \sqsubseteq, \sqcup, \sqcap)$ is a lattice over an interval domain I .

Interval (PI): $[lb, ub] \in I$, where lb is lower bound & ub is upper bound.




Interval Lattice for PIA

Formalizing Phase Interval Analysis (PIA)

Lattice : $(I, \sqsubseteq, \sqcup, \sqcap)$ is a lattice over an interval domain I .

Partial Order (\sqsubseteq): $[lb_1, ub_1] \sqsubseteq [lb_2, ub_2] \iff lb_2 \leq lb_1 \leq ub_1 \leq ub_2$.




Interval Lattice for PIA

Formalizing Phase Interval Analysis (PIA)

Lattice : $(I, \sqsubseteq, \sqcup, \sqcap)$ is a lattice over an interval domain I .

Join Operation (\sqcup): $[lb_1, ub_1] \sqcup [lb_2, ub_2] \implies [\min(lb_1, lb_2), \max(ub_1, ub_2)]$.




Interval Lattice for PIA

Formalizing Phase Interval Analysis (PIA)

Lattice : $(I, \sqsubseteq, \sqcup, \sqcap)$ is a lattice over an interval domain I .

Meet Operation (\sqcap): $[lb_1, ub_1] \sqcap [lb_2, ub_2] \implies [\max(lb_1, lb_2), \min(ub_1, ub_2)]$.



Interval Lattice for PIA

Data flow Equation for Phase Interval Analysis for a node $S \in G$ is defined as:

$$PI_{in}[S] = \bigsqcup_{p \in pred[S]} PI_{out}[p]$$

where the confluence operator is the join operation (\bigsqcup), p is an immediate predecessor node of S , and $pred[S]$ is the set of all the immediate predecessor nodes of $S \in G$.

Data flow Equation for Phase Interval Analysis for a node $S \in G$ is defined as:

$$PI_{in}[S] = \bigsqcup_{p \in pred[S]} PI_{out}[p]$$

where the confluence operator is the join operation (\bigsqcup), p is an immediate predecessor node of S , and $pred[S]$ is the set of all the immediate predecessor nodes of $S \in G$.

Transfer Function is defined as:

$$PI_{out}(S) = f(PI_{in}(S)) = \begin{cases} PI_{in}(S), & \text{Identity function, if } S \\ & \text{does not have a barrier.} \\ PI_{in}(S) + [1, 1], & \text{otherwise.} \end{cases}$$

The binary operator $+$ is defined as $[lb_1, ub_1] + [lb_2, ub_2] = [lb_1 + lb_2, ub_1 + ub_2]$.

Widening is defined as:

$$[lb_1, ub_1] \nabla_{[lb_t, ub_t]} [lb_2, ub_2] = \begin{cases} \begin{cases} [lb_t, ub_t] & \text{if } ub_1 < ub_2 \\ [lb_t, ub_1] & \text{otherwise} \end{cases} & \text{if } lb_1 < lb_2 \\ \begin{cases} [lb_1, ub_t] & \text{if } ub_1 < ub_2 \\ [lb_1, ub_1] & \text{otherwise} \end{cases} & \text{otherwise} \end{cases}$$

where lb_t and ub_t are widening thresholds.

Widening is defined as:

$$[lb_1, ub_1] \nabla_{[lbt, ubt]} [lb_2, ub_2] = \begin{cases} \begin{cases} [lbt, ubt] & \text{if } ub_1 < ub_2 \\ [lbt, ub_1] & \text{otherwise} \end{cases} & \text{if } lb_1 < lb_2 \\ \begin{cases} [lb_1, ubt] & \text{if } ub_1 < ub_2 \\ [lb_1, ub_1] & \text{otherwise} \end{cases} & \text{otherwise} \end{cases}$$

where lbt and ubt are widening thresholds.

For a loop header H and loop trip count TC , widening is applied as:

$$PI_{in}(H) = \begin{cases} PI_{in}^0(H) \sqcup (PI_{in}^0(H) + (TC * (PI_{in}^0(H) \Delta PI_{in}^1(H)))) & \text{If the trip count } TC \text{ is known} \\ PI_{in}^0(H) \nabla_{[0, \infty]} PI_{in}^1(H) & \text{otherwise} \end{cases}$$


where difference (Δ): $[lb_1, ub_1] \Delta [lb_2, ub_2] \implies [lb_2 - lb_1, ub_2 - ub_1]$.

Phase Interval Analysis: Working Example



```
1 #pragma omp parallel
2 {
3     // S1
4     for (int i = 0; i < 100; i++) {
5         // S2
6     }
7     // S3
8 }
9 // S4
10 }
```

OpenMP parallel construct




Phase Interval Analysis: Working Example



```
1 #pragma omp parallel
2 {
3     // S1
4     for (int i = 0; i < 100; i++) {
5         // S2
6     }
7     // S3
8 }
9 // S4
10 }
```

OpenMP parallel construct




Phase Interval Analysis: Working Example



```
1 #pragma omp parallel
2 {
3     // S1
4     for (int i = 0; i < 100; i++) {
5         // S2
6     }
7     // S3
8 }
9 // S4
10 }
```

OpenMP parallel construct




Phase Interval Analysis: Working Example



```
1 #pragma omp parallel
2 {
3     // S1
4     for (int i = 0; i < 100; i++) {
5         // S2
6     }
7     // S3
8 }
9 // S4
10 }
```

OpenMP parallel construct




Phase Interval Analysis: Working Example



```
1 #pragma omp parallel
2 {
3     // S1
4     for (int i = 0; i < 100; i++) {
5         // S2
6     }
7     // S3
8 }
9 // S4
10 }
```

OpenMP parallel construct




Phase Interval Analysis: Working Example



```
1 #pragma omp parallel
2 {
3     // S1
4     for (int i = 0; i < 100; i++) {
5         // S2
6     }
7     // S3
8 }
9 // S4
10 }
```

OpenMP parallel construct




Phase Interval Analysis: Working Example



```
1 #pragma omp parallel
2 {
3     // S1
4     for (int i = 0; i < 100; i++) {
5         // S2
6     }
7     // S3
8 }
9 // S4
10 }
```

OpenMP parallel construct




Phase Interval Analysis: Working Example



```
1 #pragma omp parallel
2 {
3     // S1
4     for (int i = 0; i < 100; i++) {
5         // S2
6     }
7     // S3
8 }
9 // S4
10 }
```

OpenMP parallel construct




Phase Interval Analysis: Working Example



```
1 #pragma omp parallel
2 {
3     // S1
4     for (int i = 0; i < 100; i++) {
5         // S2
6     }
7     // S3
8 }
9 // S4
10 }
```

OpenMP parallel construct




Phase Interval Analysis: Working Example



```
1 #pragma omp parallel
2 {
3     // S1
4     for (int i = 0; i < 100; i++) {
5         // S2
6     }
7     // S3
8 }
9 // S4
10 }
```

OpenMP parallel construct




Phase Interval Analysis: Working Example



```
1 #pragma omp parallel
2 {
3     // S1
4     for (int i = 0; i < 100; i++) {
5         // S2
6     }
7     // S3
8 }
9 // S4
10 }
```

OpenMP parallel construct




Phase Interval Analysis: Working Example



```
1 #pragma omp parallel
2 {
3     // S1
4     for (int i = 0; i < 100; i++) {
5         // S2
6     }
7     // S3
8 }
9 // S4
10 }
```

OpenMP parallel construct




Phase Interval Analysis: Working Example



```
1 #pragma omp parallel
2 {
3     // S1
4     for (int i = 0; i < 100; i++) {
5         // S2
6     }
7     // S3
8 }
9 // S4
10 }
```


OpenMP parallel construct



Phase Interval Analysis: Working Example

```
1 #pragma omp parallel
2 {
3     // S1
4     for (int i = 0; i < 100; i++) {
5         // S2
6     }
7     // S3
8 }
9 // S4
10 }
```


OpenMP parallel construct



Phase Interval Analysis: Working Example

```
1 #pragma omp parallel
2 {
3     // S1
4     for (int i = 0; i < 100; i++) {
5         // S2
6 #pragma omp barrier
7         // S3
8     }
9     // S4
10 }
```

OpenMP parallel construct



Technique 2: May-Happen-in-Parallel (MHP) Analysis



Given: TASKGRAPH and Phase Intervals.

Question: Can the two statements **S1** and **S2** be executed in parallel?

```
1 #pragma omp parallel shared(b, error)
2 {
3 #pragma omp for nowait
4     for(i = 0; i < len; i++)
5         a[i] = b + a[i] * 5; //S1
6 #pragma omp single
7     error = a[9] + 1; //S2
8 }
```

Technique 2: May-Happen-in-Parallel (MHP) Analysis



Given: TASKGRAPH and Phase Intervals.

Question: Can the two statements **S1** and **S2** be executed in parallel?

```
1 #pragma omp parallel shared(b, error)
2 {
3   #pragma omp for nowait
4   for(i = 0; i < len; i++)
5     a[i] = b + a[i] * 5; //S1
6   #pragma omp single
7     error = a[9] + 1; //S2
8 }
```

Answer: Instances of two nodes $u, v \in G$ may run in parallel if and only if $PI[u] \cap PI[v]$. Here the binary operator \cap on the phase intervals $PI_1 = [lb_1, ub_1]$ and $PI_2 = [lb_2, ub_2]$ is:

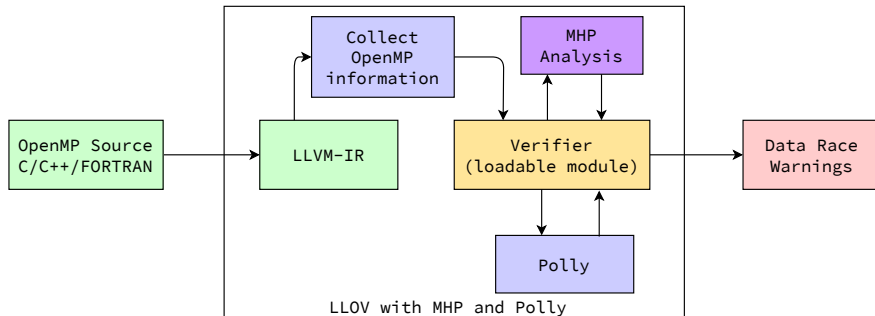
$$PI_1 \cap PI_2 = \begin{cases} true, & lb_1 \leq lb_2 \leq ub_1 \vee \\ & lb_2 \leq ub_1 \leq ub_2 \\ false, & \text{otherwise.} \end{cases}$$

Implementation Details

Designed and implemented a static **data-race** checker, **LLOV**, in the LLVM compiler framework.

LLOV:

- is based on the intermediate representation of LLVM (**LLVM-IR**)
- uses polyhedral dependence analysis
- uses May-Happen-in-Parallel analysis
- can handle OpenMP programs in C/C++/FORTRAN
- also can conservatively state when a program is **data-race free**



Flow Diagram of LLVM OpenMP Verifier (LLOV)

LLOV: Reconstructing pragmas from IR



```
1 #pragma omp parallel shared(b, error)
2 {
3 #pragma omp for nowait
4     for(i = 0; i < len; i++)
5         a[i] = b + a[i]*5;
6 #pragma omp single
7     error = a[9] + 1;
8 }
```

```
1 #pragma omp parallel shared(b, error)
2 {
3 #pragma omp for nowait
4     for(i = 0; i < len; i++)
5         a[i] = b + a[i]*5;
6 #pragma omp single
7     error = a[9] + 1;
8 }
```

Directive: OMP_Parallel

Variables:

Private: %.omp.ub = alloca i32, align 4

Private: %.omp.lb = alloca i32, align 4

Shared: i32* %i

Shared: i32* %len

Firstprivate: i64 %vla

Shared: i32* %a

Shared: i32* %b

Shared: i32* %error

Private: %.omp.stride = alloca i32, align 4

Private: %.omp.is_last = alloca i32, align 4

Child Directives:

1: Directive: OMP_Workshare_Loop

Schedule type : Static Schedule (auto-chunked)

2: Directive: OMP_Workshare_single

3: Directive: OMP_Barrier

In-memory representation of a directive.

```
1 #pragma omp parallel shared(b, error)
2 {
3 #pragma omp for nowait
4     for(i = 0; i < len; i++)
5         a[i] = b + a[i]*5;
6 #pragma omp single
7     error = a[9] + 1;
8 }
```

Directive: OMP_Parallel

Variables:

Private: %.omp.ub = alloca i32, align 4

Private: %.omp.lb = alloca i32, align 4

Shared: i32* %i

Shared: i32* %len

Firstprivate: i64 %vla

Shared: i32* %a

Shared: i32* %b

Shared: i32* %error

Private: %.omp.stride = alloca i32, align 4

Private: %.omp.is_last = alloca i32, align 4

Child Directives:

1: Directive: OMP_Workshare_Loop

Schedule type : Static Schedule (auto-chunked)

2: Directive: OMP_Workshare_single

3: Directive: OMP_Barrier

<Directive> ::= <Dtype> [<Sched>] { <Var> } { <Directive> }

<Dtype> ::= parallel | for | simd
| workshare | sections
| section | single | master
| critical | atomic
| target | distribute

<Var> ::= <Vtype> val

<Vtype> ::= private | firstprivate
| lastprivate | threadprivate
| shared | reduction

<Sched> ::= [<modifier>] [ordered] <Stype> <chunk>

<modifier> ::= monotonic | nonmonotonic

<Stype> ::= static | dynamic
| guided | auto | runtime

<chunk> ::= <non-negative integer>

In-memory representation of a directive.

Data-Race: The **source memory access** MA_1 (in basic block u) and **sink memory access** MA_2 (in basic block v) of a data dependence may potentially **race** if $PI[u] \cap PI[v]$.

Data Race detected.

Source : llvm/lib/Transforms/OpenMPVerify/test/10.race1.c:10:11

Sink : llvm/lib/Transforms/OpenMPVerify/test/10.race1.c:12:11

=====

```
9 :      {  
10 :    x = 1;  
11 : #pragma omp section  
12 :    x = 2;  
13 :      }
```

=====

Error reporting: LLOV highlights the source location of the race condition.

Experimental Evaluation

We evaluate our work on two metrics:

- Compare **Performance** with seeded benchmarks
- Measure **Scalability** with real world applications

Benchmarks

- DataRaceBench C/C++ v1.3.2 [Verma et al., 2020] to evaluate performance
 - 172 Kernels, 83 kernels with data-races, 89 kernels without any data-race
- 22 ECP Proxy application to evaluate scalability (1 million LOC)
- DataRaceBench FORTRAN [Kukreja et al., 2019]

Benchmarks

- DataRaceBench C/C++ v1.3.2 [Verma et al., 2020] to evaluate performance
 - 172 Kernels, 83 kernels with data-races, 89 kernels without any data-race
- 22 ECP Proxy application to evaluate scalability (1 million LOC)
- DataRaceBench FORTRAN [Kukreja et al., 2019]

System Specifications

System: Two Intel Xeon E5-2697 v4 @ 2.30GHz processors

OS: 64 bit Ubuntu 20.04.2 LTS Server

Kernel: Linux kernel version 5.4.0-71-generic

Threads: 72 (2 x 36) hardware threads

Memory: 128GB


OpenMP library: LLVM OpenMP runtime v5.0.1 (libomp5)

Race detection tools with the version numbers and flags used for comparison. [Dynamic tool](#), [Static tool](#).

Tools	Version	Flags	OpenMP	Compiler
HELGRIND [Valgrind-project, 2007b]	3.15.0	<code>-tool=helgrind</code>	gomp	gcc 9.3
VALGRIND DRD [Valgrind-project, 2007a]	3.15.0	<code>-tool=drd -check-stack-var=yes</code>	gomp	gcc 9.3
TSAN-LLVM [Serebryany et al., 2011]	12.0.1	<code>ignore_noninstrumented_modules=1</code>	lomp	clang 12.0
ARCHER [Atzeni et al., 2016]	5b37681	<code>ignore_noninstrumented_modules=1</code>	lomp	clang 6.0
SWORD [Atzeni et al., 2018]	7a08f3c	<code>-analysis-tool=sword-race-analysis</code>	lomp	clang 7.1
ROMP [Gu and Mellor-Crummey, 2018]	v2.0		gomp	gcc 9.3
OMPRACER [Swain et al., 2020]	0.1.1	<code>-no-filter -silent -nolimit</code>	lomp	clang 9.0
LLOV (WITH POLLY)	v0.3	<code>openmp-verify-use-llvm-da=false</code>	lomp	clang 12.0
LLOV (WITH MHP)	v0.3	<code>openmp-verify-use-polly-da=false</code>	lomp	clang 12.0
LLOV (WITH MHP & POLLY)	v0.3	<code>-Xclang -disable-O0-optnone</code>	lomp	clang 12.0

Unavailable tools: [OMPVERIFY](#) [Basupalli et al., 2011], [POLYOMP](#) [Chatarasi et al., 2016], and [DRACO](#) [Ye et al., 2018].

Performance: DataRaceBench v1.3.2 Results



Tools (nt for without tasks, p for Polly, m for MHP, mp for MHP and Polly)

Maximum number of races reported by different tools in DataRaceBench v1.3.2


Precision, Recall, and Accuracy of the tools on DataRaceBench v1.3.2

Tools	Precision	Recall	Accuracy	F1 Score	Diagnostic Odds Ratio
HELGRIND	0.50	0.94	0.51	0.65	1.72
VALGRIND DRD	0.61	0.94	0.68	0.74	13.16
TSAN-LLVM	0.66	0.93	0.73	0.77	15.58
ARCHER	0.91	0.85	0.89	0.88	64.07
SWORD	0.89	0.73	0.79	0.80	20.13
ROMP	0.85	0.79	0.83	0.82	24.93
OMPRACER	0.85	0.81	0.83	0.83	25.30
OMPRACER (notasks)	0.86	0.92	0.87	0.88	48.26
LLOV (WITH POLLY)	0.94	0.87	0.88	0.90	52.57
LLOV (WITH MHP)	0.74	0.99	0.82	0.85	131.25
LLOV (WITH MHP & POLLY)	0.90	0.99	0.94	0.94	525.00

22 Miscellaneous Benchmarks used for Scalability Analysis


Benchmark	Domain	KLOC	Benchmark	Domain	KLOC
AMG	Parallel Algebraic Multigrid Solver	64.9	NBody	Parallel N-body Simulation	0.4
CoMD	Molecular Dynamics Simulation	6.1	openmp-tutorial	SC OpenMP Tutorial	10.4
COVID-SIM	COVID-19 Transmission Dynamics	19.0	OpenMP-Microbench	OpenMP Performance Benchmarking	1.0
LCALS	Livermore Loops	6.6	ParallelResearchKernels	Parallel Architectures	106.2
LULESH	Shock Hydrodynamics	5.5	QuickSilver	Monte Carlo Particle Transport	10.0
miniAMR	Adaptive Mesh Refinement	17.7	Rodinia	Parallel Computing Benchmark Suite	186.4
miniBUDE	Molecular Docking	9.5	RSBench	Monte Carlo Neutron Transport	6.0
miniFE	Finite Element	297.8	sw4lite	Geodynamics	54.2
miniQMC	Quantum Monte Carlo Simulation	31.9	Sundials	Nonlinear ODE Solver	186.9
miniVite	Louvain Method for Community Detection	2.4	TriangleCounting	Parallel Triangle Counting	0.8
NAS-Parallel	Performance of Parallel Supercomputers	15.9	XSBench	Monte Carlo Neutron Transport	6.7

Combined total of 1 million lines of code.



Cactus plot of different race detection tools on 22 miscellaneous benchmarks.

Scalability Analysis: Memory Footprint (Goal 3)



Average maximum resident set size (RSS) of the tools across all the 22 miscellaneous benchmarks

An implementation of DataRaceBench C/C++ v1.2 [Liao et al., 2017] in FORTRAN 95.

- Demonstrate LLOV is language agnostic
- Created DataRaceBench FORTRAN
 - Converted 92 (out of 116) C/C++ kernels
- Open-sourced the benchmark [Kukreja et al., 2019]

¹Earlier version of LLOV. FORTRAN frontend (FLANG) generating LLVM-IR is replaced by a new frontend (F18) that does not yet generate LLVM-IR.

An implementation of DataRaceBench C/C++ v1.2 [Liao et al., 2017] in FORTRAN 95.

- Demonstrate LLOV is language agnostic
- Created DataRaceBench FORTRAN
 - Converted 92 (out of 116) C/C++ kernels
- Open-sourced the benchmark [Kukreja et al., 2019]

Maximum number of races reported by different tools in DataRaceBench FORTRAN

Tools	Race: Yes		Race: No		Coverage/92
	TP	FN	TN	FP	
HELGRIND	46	6	4	36	92
VALGRIND DRD	45	7	21	19	92
LLOV v0.2 ¹	36	7	19	5	67

¹Earlier version of LLOV. FORTRAN frontend (FLANG) generating LLVM-IR is replaced by a new frontend (F18) that does not yet generate LLVM-IR.

Comparison of OpenMP pragma handling by OpenMP aware tools.(Y for Yes, N for No)

OpenMP Pragma	LLOV	OMPRACER	OMPVERIFY	POLYOMP	DRACO	SWORD	ARCHER	ROMP
#pragma omp parallel	Y	Y	Y	Y	Y	Y	Y	Y
#pragma omp for	Y	Y	Y	Y	Y	Y	Y	Y
#pragma omp parallel for	Y	Y	Y	Y	Y	Y	Y	Y
#pragma omp critical	Y	Y	N	N	N	Y	Y	Y
#pragma omp atomic	Y	Y	N	N	N	Y	Y	Y
#pragma omp master	Y	Y	N	Y	N	Y	Y	Y
#pragma omp single	Y	Y	N	Y	N	Y	Y	Y
#pragma omp simd	Y	Y	N	N	Y	N	N	N
#pragma omp parallel for simd	Y	Y	N	N	Y	N	N	N
#pragma omp parallel sections	Y	N	N	N	N	Y	Y	Y
#pragma omp sections	Y	N	N	N	N	Y	Y	Y
#pragma omp threadprivate	Y	Y	N	N	N	N	Y	Y
#pragma omp ordered	Y	N	N	N	N	N	Y	Y
#pragma omp distribute	Y	Y	N	N	N	N	Y	Y
#pragma omp task	N	N	N	N	N	N	Y	Y
#pragma omp taskgroup	N	N	N	N	N	N	Y	Y
#pragma omp taskloop	N	N	N	N	N	N	Y	Y
#pragma omp taskwait	N	N	N	N	N	N	Y	Y
#pragma omp barrier	Y	Y	N	Y	N	Y	Y	Y
#pragma omp teams	Y	Y	N	N	N	N	N	N
#pragma omp target	Y	Y	N	N	N	N	N	N

We proposed and implemented a fast, static **data-race** checker for shared-memory parallel programs using

- Polyhedral Dependence Analysis
- May-Happen-in-Parallel (MHP) Analysis

We proposed and implemented a fast, static **data-race** checker for shared-memory parallel programs using

- Polyhedral Dependence Analysis
- May-Happen-in-Parallel (MHP) Analysis

Observations

- Using only precise dependence analysis results in **low coverage** but **high precision**
- Using only MHP analysis results in **more coverage** but **less precision**

Conclusion

- Combination of precise dependence analysis and MHP analysis results in **high precision**, **high accuracy**, and **better coverage**

Handling explicit tasks in OpenMP

Support tasking constructs: task, taskloop, taskgroup, and taskwait.
None of the static techniques support tasking. (Open Problem).

MLIR based LLOV

Use OpenMP dialect of MLIR in place of LLVM-IR.

PIA for Optimization

PIA for C++ threads and pthreads will be interesting, and will expose classic compiler optimizations on shared-memory parallel programs.

- **Utpal Bora**, Santanu Das, Pankaj Kukreja, Saurabh Joshi, Ramakrishna Upadrasta, and Sanjay Rajopadhye. 2020. “LLOV: A Fast Static Data-Race Checker for OpenMP Programs.” *ACM Trans. Archit. Code Optim.* 17, 4, Article 35 (December 2020), 26 pages. DOI:<https://doi.org/10.1145/3418597>
- **Utpal Bora**, Shraiys Vaishay, Saurabh Joshi and Ramakrishna Upadrasta, “OpenMP aware MHP Analysis for Improved Static Data-Race Detection”, *2021 IEEE/ACM 7th Workshop on the LLVM Compiler Infrastructure in HPC (LLVM-HPC)*, 2021, pp. 1-11, doi: 10.1109/LLVMHPC54804.2021.00006
- Shalini Jain, **Utpal Bora**, Prateek Kumar, Vaibhav B. Sinha, Suresh Purini and Ramakrishna Upadrasta, “An analysis of executable size reduction by LLVM passes.” *CSI Transactions on ICT* 7, 105110 (2019). doi: 10.1007/s40012-019-00248-5

Thank you for your time! Questions?

LLOV is freely available for download.

Link: <https://github.com/utpalbora/llov>

Blog: <https://compilers.cse.iith.ac.in/projects/llov/>



utpalbora.com
(ub230@cam.ac.uk)

Thank You!!



Atzeni, S., Gopalakrishnan, G., Rakamarić, Z., Ahn, D. H., Laguna, I., Schulz, M., Lee, G. L., Protze, J., and Müller, M. S. (2016).

ARCHER: effectively spotting data races in large openmp applications.

In *2016 IEEE International Parallel and Distributed Processing Symposium, IPDPS 2016, Chicago, IL, USA, May 23-27, 2016*, pages 53–62. IEEE Computer Society.



Atzeni, S., Gopalakrishnan, G., Rakamarić, Z., Laguna, I., Lee, G. L., and Ahn, D. H. (2018).

SWORD: A bounded memory-overhead detector of openmp data races in production runs.

In *2018 IEEE International Parallel and Distributed Processing Symposium, IPDPS 2018, Vancouver, BC, Canada, May 21-25, 2018*, pages 845–854. IEEE Computer Society.



Basupalli, V., Yuki, T., Rajopadhye, S., Morvan, A., Derrien, S., Quinton, P., and Wonnacott, D. (2011).

ompverify: Polyhedral analysis for the openmp programmer.

In Chapman, B. M., Gropp, W. D., Kumaran, K., and Müller, M. S., editors, *OpenMP in the Petascale Era - 7th International Workshop on OpenMP, IWOMP 2011, Chicago, IL, USA, June 13-15, 2011. Proceedings*, volume 6665 of *Lecture Notes in Computer Science*, pages 37–53, Berlin, Heidelberg. Springer Berlin Heidelberg.



Chatarasi, P., Shirako, J., Kong, M., and Sarkar, V. (2016).

An extended polyhedral model for SPMD programs and its use in static data race detection.

In Ding, C., Criswell, J., and Wu, P., editors, *Languages and Compilers for Parallel Computing - 29th International Workshop, LCPC 2016, Rochester, NY, USA, September 28-30, 2016, Revised Papers*, volume 10136 of *Lecture Notes in Computer Science*, pages 106–120. Springer.



Gu, Y. and Mellor-Crummey, J. (2018).

Dynamic data race detection for openmp programs.

In *Proceedings of the International Conference for High Performance Computing, Networking, Storage, and Analysis, SC '18*. IEEE Press.



Kildall, G. A. (1973).

A unified approach to global program optimization.

In *Proceedings of the 1st Annual ACM SIGACT-SIGPLAN Symposium on Principles of Programming Languages*, POPL '73, page 194206, New York, NY, USA. Association for Computing Machinery.



Kukreja, P., Shukla, H., and Bora, U. (2019).

DataRaceBench FORTRAN.

https://github.com/IITH-Compilers/dr_b_fortran.

[Online; accessed 19-October-2019].



Liao, C., Lin, P.-H., Asplund, J., Schordan, M., and Karlin, I. (2017).

Dataracebench: A benchmark suite for systematic evaluation of data race detection tools.

In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*, SC '17, pages 11:1–11:14, New York, NY, USA. ACM.



Mayer, F., Knaust, M., and Philippsen, M. (2019).

Openmp on fpgas—a survey.

In Fan, X., de Supinski, B. R., Sinnen, O., and Giacaman, N., editors, *OpenMP: Conquering the Full Hardware Spectrum*, pages 94–108, Cham. Springer International Publishing.



Serebryany, K., Potapenko, A., Iskhodzhanov, T., and Vyukov, D. (2011).

Dynamic race detection with llvm compiler.

In *Proceedings of the Second International Conference on Runtime Verification*, RV'11, pages 110–114, Berlin, Heidelberg. Springer-Verlag.



Swain, B., Li, Y., Liu, P., Laguna, I., Georgakoudis, G., and Huang, J. (2020).

OMPRacer: A Scalable and Precise Static Race Detector for OpenMP Programs.

In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*, SC '20. IEEE Press.



Valgrind-project (2007a).

DRD: a thread error detector.

<http://valgrind.org/docs/manual/drd-manual.html>.

[Online; accessed 08-May-2019].



Valgrind-project (2007b).

Helgrind: a thread error detector.

<http://valgrind.org/docs/manual/hg-manual.html>.

[Online; accessed 08-May-2019].



Verma, G., Shi, Y., Liao, C., Chapman, B., and Yan, Y. (2020).

Enhancing dataracebench for evaluating data race detection tools.

In *2020 IEEE/ACM 4th International Workshop on Software Correctness for HPC Applications (Correctness)*, pages 20–30.



Ye, F., Schordan, M., Liao, C., Lin, P., Karlin, I., and Sarkar, V. (2018).

Using polyhedral analysis to verify openmp applications are data race free.

In Laguna, I. and Rubio-González, C., editors, *2nd IEEE/ACM International Workshop on Software Correctness for HPC Applications, CORRECTNESS@SC 2018, Dallas, TX, USA, November 12, 2018*, pages 42–50. IEEE.

Extra Slides!

- Implemented a *fast, static, and language agnostic* OpenMP **data-race** checker
- Demonstrated use of Polyhedral dependences and MHP in **data-race** detection
- Proposed and Formalized OpenMP aware MHP Analysis using Phase Interval Analysis
- Created DataRaceBench FORTRAN, an annotated benchmark of **data-races** in FORTRAN
- Visualization of TASKGRAPH for better developer productivity

Data-Races in Miscellaneous Benchmarks



Number of Data-Races reported by various tools in 22 Miscellaneous Benchmarks

Benchmark	KLOC	LLOV	TSAN	ARCHER	OMPRACER	ROMP	HELGRIND	DRD
AMG	64.9	644	536	97	2489	9520	1000	1000
CoMD	6.1	22	71	0	12	0	367	129
COVID-SIM	19.0	331	2	0	0	0	350	0
LCALS	6.6	18	0	0	0	0	0	0
LULESH	5.5	126	71	0	139	0	1000	648
miniAMR	17.7	7693	0	250	539	146644	159	106
miniBUDE	9.5	0	1	0	0	0	131	57
miniFE	297.8	0	177	0	2	0	743	332
miniQMC	31.9	593	90	0	0	0	193	143
miniVite	2.4	5	66	0	68	14048	393	218
NAS-Parallel	15.9	219	17	84	8	0	73	10
NBody	0.4	18	99	0	4	0	133	140
openmp-tutorial	10.4	4	462	1	1	0	203	24
OpenMP-Microbench	1.0	0	46	2	0	0	335	19
ParallelResearchKernels	106.2	31	13	31	4165	36840	105	33
QuickSilver	10.0	0	48	33	7	0	128	2
Rodinia	186.4	55	166	3590	33	2231871	206	121
RSBench	6.0	0	4	0	0	0	13	0
sw4lite	54.2	42	895	0	0	0	112	112
Sundials	186.9	59	13	64	0	0	164	78
TriangleCounting	0.8	11	100	1	0	0	576	273
XSBench	6.7	0	6	0	0	0	23	0

```
1: Input: TASKGRAPH  $G$ 
2: Output: Phase Intervals  $PI$ 
3: Data: worklist
4: function PIA( $G$ )
5:   worklist.add( $G.root$ )
6:   for all  $n \in G$  do
7:      $PI_{out}[n] \leftarrow \perp$  ▷ init to bot
8:     if  $n$  has barrier then worklist.add( $n$ ) end if
9:   end for
10:   $PI_{in}[G.root] \leftarrow [0, 0]$ 
11:  while !worklist.empty() do
12:     $n \leftarrow$  worklist.pop()
13:     $PI_{in}[n] \leftarrow \bigsqcup_{p \in pred[n]} PI_{out}[p]$  ▷ join
```

```
14:      if  $n$  is loop header then
15:           $Pl_{out}[n] \leftarrow w(Pl_{in}[n])$                                 ▷ widen
16:      end if
17:       $Pl_{out}[n] \leftarrow f(Pl_{in}[n])$                                 ▷ transfer
18:      if  $Pl_{out}[n] \neq Pl_{in}[n]$  then
19:          for all  $s \in \text{successor}(n)$  do
20:               $\text{worklist.add}(s)$ 
21:          end for
22:      end if
23:  end while
24:  return  $Pl_{out}$ 
25: end function
```

Race detection literature (minimalistic)

Published	Title	Year
TOCS	Eraser: A dynamic data race detector for multithreaded programs	1997
SOSP	RacerX: Effective, Static Detection of Race Conditions and Deadlocks	2003
POPL	Conditional must not aliasing for static race detection	2007
PPoPP	May-happen-in-parallel analysis of X10 program	2007
ESEC-FSE	RELAY: static race detection on millions of lines of code	2007
WBI	ThreadSanitizer data race detection in practice	2009
IWOMP	ompVerify : Polyhedral Analysis for the OpenMP Programmer	2011
TOPLAS	LOCKSMITH: Practical static race detection for C	2011
LCPC	An extended polyhedral model for SPMD programs and its use in static data race detection	2016
IPDPS	ARCHER: effectively spotting data races in large OpenMP applications	2016
IPDPS	SWORD: A Bounded Memory-Overhead Detector of OpenMP Data Races in Production Runs	2018
Correctness	Using Polyhedral Analysis to Verify OpenMP Applications are Data Race Free	2018
SC	Dynamic Data Race Detection for OpenMP Programs	2018
OOPSLA	RacerD: Compositional Static Race Detection	2018
IWOMP	OMPSan: Static Verification of OpenMPs Data Mapping Constructs	2019
SC	OMPRacer: A Scalable and Precise Static Race Detector for OpenMP Programs	2020
TACO	LLOV: A Fast Static Data-Race Checker for OpenMP Programs	2020

The only FN case for LLOV



```
1  #pragma omp parallel shared(a) private(i)
2  {
3      #pragma omp master
4      a = 0;
5
6      #pragma omp for reduction(+:a)
7      for (i=0; i<10; i++){
8          a = a + i;
9      }
10
11     #pragma omp single
12     printf("Sum is %d\n", a);
13 }
```

DRB140: Example of the only FN case for LLOV

Improvement area: PIA can be made more precise



```
1  #pragma omp parallel
2      for(int i=0; i<N ;i++){
3          #pragma omp for
4          for(int i=0; i<C; i++){
5              temp[i] = b[i] + c[i];
6          }
7
8          #pragma omp for
9          for(int i=C-1; i>=0; i--){
10             b[i] = temp[i] * a;
11         }
12     }
```

DRB159: PIA is not precise due to conservative analysis

- **True Positive (TP):** If the evaluation tool correctly detects a data-race present in the kernel it is a True Positive test result. A higher number of true positives represents a better tool.
- **True Negative (TN):** If the benchmark does not contain a race and the tool declares it as race-free, then it is a true negative case. A higher number of true negatives represents a better tool.
- **False Positives (FP):** If the benchmark does not contain any race, but the tool reports a race condition, it is a false positive. False Positives should be as low as possible.
- **False Negatives (FN):** False Negative test result is obtained when the tool fails to detect a known race in the benchmark. These are the cases that are missed by the tool. A lower number of false negatives are desirable.

- **Precision** : Precision is the measure of closeness of the outcomes of prediction. Thus, a higher value of precision represents that the tool will more often than not identify a race condition when it exists.

$$Precision = \frac{TP}{TP + FP}$$

- **Recall** : Recall gives the total number of cases detected out of the maximum data-races present. A higher recall value means that there are less chances that a data-race is missed by the tool. It is also called true positive rate (TPR).

$$Recall = \frac{TP}{TP + FN}$$

- **Accuracy** : Accuracy gives the chances of correct reports out of all the reports, as the name suggests. A higher value of accuracy is always desired and gives overall measure of the efficacy of the tool.

$$Accuracy = \frac{TP + TN}{TP + FP + TN + FN}$$

- **F1 Score** : The harmonic mean of precision and recall is called the F1 score. An F1 score of 1 can be achieved in the best case when both precision and recall are perfect. The worst case F1 score is 0 when either precision or recall is 0.

$$F1\ Score = 2 * \frac{Precision * Recall}{Precision + Recall}$$

- **Diagnostic odds ratio (DOR)** : It is the ratio of the positive likelihood ratio (LR+) to the negative likelihood ratio (LR-).

$$DOR = \frac{LR+}{LR-} \text{ where,}$$

$$\text{Positive Likelihood Ratio (LR+)} = \frac{TPR}{FPR},$$

$$\text{Negative Likelihood Ratio (LR-)} = \frac{FNR}{TNR},$$

$$\text{True Positive Rate (TPR)} = \frac{TP}{TP + FN},$$

$$\text{False Positive Rate (FPR)} = \frac{FP}{FP + TN},$$

$$\text{False Negative Rate (FNR)} = \frac{FN}{FN + TP} \text{ and}$$

$$\text{True Negative Rate (TNR)} = \frac{TN}{TN + FP}$$

DOR is the measure of the ratio of the odds of race detection being positive given that the test case has a data-race, to the odds of race detection being positive given the test case does not have a race.